

Dates and Times with lubridate

In this lesson, we'll explore the `lubridate` R package, by **Garrett Golemund** and **Hadley Wickham**. According to the package authors, *"lubridate has a consistent, memorable syntax, that makes working with dates fun instead of frustrating."* If you've ever worked with dates in R, that statement probably has your attention.

Unfortunately, due to different date and time representations, this lesson is only guaranteed to work with an **"en_US.UTF-8"** locale. To view your locale, type `Sys.getlocale("LC_TIME")`.

```
Sys.getlocale("LC_TIME")
```

```
[1] "en_NZ.UTF-8"
```

If the output above is not **"en_US.UTF-8"**, this lesson is not guaranteed to work correctly. Of course, you are welcome to try it anyway. We apologize for this inconvenience.

`lubridate` was automatically installed (if necessary) and loaded upon starting this lesson. To build the habit, we'll go ahead and (re)load the package now. Type `library(lubridate)` to do so.

```
library(lubridate)
```

`lubridate` contains many useful functions. We'll only be covering the basics here. Type `help(package = lubridate)` to bring up an overview of the package, including the package DESCRIPTION, a list of available functions, and a link to the official package vignette.

```
help(package = lubridate)
```

The **today()** function returns today's date. Give it a try, storing the result in a new variable called *this_day*.

```
this_day <- today()
```

Print the contents of `this_day` to the console.

```
this_day
```

```
[1] "2018-07-19"
```

There are three components to this date. In order, they are year, month, and day. We can extract any of these components using the **year()**, **month()**, or **day()** function, respectively. Try any of those on *this_day* now.

```
day(this_day)
```

```
[1] 18
```

We can also get the day of the week from *this_day* using the **wday()** function. It will be represented as a number, such that 1 = Sunday, 2 = Monday, 3 = Tuesday, etc. Give it a shot.

```
wday(this_day)
```

```
[1] 4
```

Now try the same thing again, except this time add a second argument, *label = TRUE*, to display the *name* of the weekday (represented as an ordered factor).

```
wday(this_day, label = TRUE)
```

```
[1] Wed
Levels: Sun < Mon < Tue < Wed < Thu < Fri < Sat
```

In addition to handling dates, `lubridate` is great for working with date and time combinations, referred to as date-times. The **now()** function returns the date-time representing this exact moment in time. Give it a try and store the result in a variable called `this_moment`.

```
this_moment <- now()
```

```
this_moment
```

```
[1] "2018-07-19 16:02:22 NZST"
```

Just like with dates, we can extract the *year*, *month*, *day*, or *day of week*. However, we can also use **hour()**, **minute()**, and **second()** to extract specific time information. Try any of these three new functions now to extract one piece of time information from *this_moment*.

```
second(this_moment)
```

```
[1] 22.54331
```

today() and **now()** provide neatly formatted date-time information. When working with dates and times ‘in the wild’, this won’t always (and perhaps rarely will) be the case.

Fortunately, lubridate offers a variety of functions for parsing date-times. These functions take the form of **ymd()**, **dmy()**, **hms()**, **ymd_hms()**, etc., where each letter in the name of the function stands for the location of *years (y)*, *months (m)*, *days (d)*, *hours (h)*, *minutes (m)*, and/or *seconds (s)* in the date-time being read in.

To see how these functions work, try `ymd(“1989-05-17”)`. You must surround the date with quotes. Store the result in a variable called *my_date*.

```
my_date <- ymd("1989-05-17")
```

Now take a look at *my_date*.

```
my_date
```

```
[1] "1989-05-17"
```

It looks almost the same, except for the addition of a time zone, which we’ll discuss later in the lesson. Below the surface, there’s another important change that takes place when lubridate parses a date. Type `class(my_date)` to see what that is.

```
class(my_date)
```

```
[1] "Date"
```

So **ymd()** took a character string as input and returned an object of class *POSIXct*. It’s not necessary that you understand what *POSIXct* is, but just know that it is one way that R stores date-time information internally.

“1989-05-17” is a fairly standard format, but lubridate is ‘smart’ enough to figure out many different date-time formats. Use **ymd()** to parse “1989 May 17”. Don’t forget to put quotes around the date!

```
ymd("1989 May 17")
```

```
[1] "1989-05-17"
```

Despite being formatted differently, the last two dates had the year first, then the month, then the day. Hence, we used **ymd()** to parse them. What do you think the appropriate function is for parsing “March 12, 1975”? Give it a try.

```
mdy("March 12, 1975")
```

```
[1] "1975-03-12"
```

We can even throw something funky at it and lubridate will often know the right thing to do. Parse 25081985, which is supposed to represent the 25th day of August 1985. Note that we are actually parsing a numeric value here – not a character string – so leave off the quotes.

```
dmy(25081985)
```

```
[1] "1985-08-25"
```

But be careful, it’s not magic. Try `ymd(“192012”)` to see what happens when we give it something more ambiguous. Surround the number with quotes again, just to be consistent with the way most dates are represented (as character strings).

```
ymd("192012")
```

```
All formats failed to parse. No formats found.
```

```
[1] NA
```

You got a warning message because it was unclear what date you wanted. When in doubt, it’s best to be more explicit. Repeat the same command, but add two dashes OR two forward slashes to “192012” so that it’s clear we want January 2, 1920.

```
ymd("1920-1-2")
```

```
[1] "1920-01-02"
```

In addition to dates, we can parse date-times. I've created a date-time object called **dt1**. Take a look at it now.

dt1

Hide

[1] "2014-08-23 17:23:02"

Now parse *dt1* with **ymd_hms()**.

ymd_hms(dt1)

Hide

[1] "2014-08-23 17:23:02 UTC"

What if we have a time, but no date? Use the appropriate lubridate function to parse “03:22:14” (hh:mm:ss).

hms("03:22:14")

Hide

[1] "3H 22M 14S"

lubridate is also capable of handling vectors of dates, which is particularly helpful when you need to parse an entire column of data. I've created a vector of dates called **dt2**. View its contents now.

dt2

Hide

[1] "2014-05-14" "2014-09-22" "2014-07-11"

Now parse *dt2* using the appropriate lubridate function.

ymd(dt2)

Hide

[1] "2014-05-14" "2014-09-22" "2014-07-11"

The **update()** function allows us to update one or more components of a date-time. For example, let's say the current time is 08:34:55 (hh:mm:ss). Update *this_moment* to the new time using the following command: *update(this_moment, hours = 8, minutes = 34, seconds = 55)*.

update(this_moment, hours = 8, minutes = 34, seconds = 55)

Hide

[1] "2018-07-18 08:34:55 NZST"

It's important to recognize that the previous command does not alter *this_moment* unless we reassign the result to *this_moment*. To see this, print the contents of *this_moment*.

this_moment

Hide

[1] "2018-07-18 13:33:10 NZST"

Unless you're a superhero, some time has passed since you first created *this_moment*. Use **update()** to make it match the current time, specifying at least hours and minutes. Assign the result to *this_moment*, so that *this_moment* will contain the new time.

this_moment <- update(this_moment, hours = 13, minutes = 33)

Hide

Take one more look at *this_moment* to see that it's been updated.

this_moment

Hide

[1] "2018-07-18 13:33:10 NZST"

Now, pretend you are in New York City and you are planning to visit a friend in Hong Kong. You seem to have misplaced your itinerary, but you know that your flight departs New York at 17:34 (5:34pm) the day after tomorrow. You also know that your flight is scheduled to arrive in Hong Kong exactly 15 hours and 50 minutes after departure.

Let's reconstruct your itinerary from what you can remember, starting with the full date and time of your departure. We will approach this by finding the current date in New York, adding 2 full days, then setting the time to 17:34.

To find the current date in New York, we'll use the **now()** function again. This time, however, we'll specify the time zone that we want: “America/New_York”. Store the result in a variable called *nyc*. Check out *?now* if you need help.

nyc <- now(tzone = "America/New_York")

Hide

For a complete list of valid time zones for use with lubridate, check out the following Wikipedia page: http://en.wikipedia.org/wiki/List_of_tz_database_time_zones

View the contents of *nyc*, which now contains the current date and time in New York.

```
nyc
```

Hide

```
[1] "2018-07-18 23:57:02 EDT"
```

Your flight is the day after tomorrow (in New York time), so we want to add two days to *nyc*. One nice aspect of *lubridate* is that it allows you to use arithmetic operators on dates and times. In this case, we'd like to add two days to *nyc*, so we can use the following expression: *nyc* + *days*(2). Give it a try, storing the result in a variable called *depart*.

```
depart <- nyc + days(2)
```

Hide

Take a look at *depart*.

```
depart
```

Hide

```
[1] "2018-07-20 23:57:02 EDT"
```

So now *depart* contains the date of the day after tomorrow. Use **update()** to add the correct hours (17) and minutes (34) to *depart*. Reassign the result to *depart*.

```
depart <- update(depart, hours = 17, minutes = 34)
```

Hide

Take one more look at *depart*.

```
depart
```

Hide

```
[1] "2018-07-20 17:34:02 EDT"
```

Your friend wants to know what time she should pick you up from the airport in Hong Kong. Now that we have the exact date and time of your departure from New York, we can figure out the exact time of your arrival in Hong Kong.

The first step is to add 15 hours and 50 minutes to your departure time. Recall that *nyc* + *days*(2) added two days to the current time in New York. Use the same approach to add 15 hours and 50 minutes to the date-time stored in *depart*. Store the result in a new variable called *arrive*.

```
arrive <- depart + hours(15) + minutes(50)
```

Hide

The *arrive* variable contains the time that it will be in New York when you arrive in Hong Kong. What we really want to know is what time it will be in Hong Kong when you arrive, so that your friend knows when to meet you.

The **with_tz()** function returns a date-time as it would appear in another time zone. Use *?with_tz* to check out the documentation.

```
?with_tz
```

Hide

Description: *with_tz* returns a date-time as it would appear in a different time zone. The actual moment of time measured does not change, just the time zone it is measured in. *with_tz* defaults to the Universal Coordinated time zone (UTC) when an unrecognized time zone is inputted. See **Sys.timezone()** for more information on how R recognizes time zones. **Usage:** *with_tz*(time, tzone = "")

Use **with_tz()** to convert *arrive* to the “Asia/Hong_Kong” time zone. Reassign the result to *arrive*, so that it will get the new value.

```
arrive <- with_tz(arrive, tzone = "Asia/Hong_Kong")
```

Hide

Print the value of *arrive* to the console, so that you can tell your friend what time to pick you up from the airport.

```
arrive
```

Hide

```
[1] "2018-07-21 21:24:02 HKT"
```

Fast forward to your arrival in Hong Kong. You and your friend have just met at the airport and you realize that the last time you were together was in Singapore on *June 17, 2008*. Naturally, you'd like to know exactly how long it has been.

Use the appropriate *lubridate* function to parse “June 17, 2008”, just like you did near the beginning of this lesson. This time, however, you should specify an extra argument, *tz* = “Singapore”. Store the result in a variable called *last_time*.

```
last_time <- mdy("June 17, 2008", tz = "Singapore")
```

Hide

View the contents of *last_time*.

```
last_time
```

Hide

```
[1] "2008-06-17 +08"
```

Pull up the documentation for **interval()**, which we'll use to explore how much time has passed between *arrive* and *last_time*.

```
?interval
```

Hide

Description: *creates an Interval object with the specified start and end dates. If the start date occurs before the end date, the interval will be positive. Otherwise, it will be negative.* **Usage:** *interval(start, end = NULL, tzone = tz(start))*

Create an **interval()** that spans from *last_time* to *arrive*. Store it in a new variable called *how_long*.

```
how_long <- interval(last_time, arrive)
```

Hide

Now use *as.period(how_long)* to see how long it's been.

```
as.period(how_long)
```

Hide

```
[1] "10y 1m 4d 21H 24M 2S"
```

This is where things get a little tricky. Because of things like leap years, leap seconds, and daylight savings time, the length of any given minute, day, month, week, or year is relative to when it occurs. In contrast, the length of a second is always the same, regardless of when it occurs.

To address these complexities, the authors of lubridate introduce four classes of time related objects: *instants*, *intervals*, *durations*, and *periods*. These topics are beyond the scope of this lesson, but you can find a complete discussion in the 2011 Journal of Statistical Software paper titled ‘**Dates and Times Made Easy with lubridate**’.

This concludes our introduction to working with dates and times in lubridate. I created a little timer that started running in the background when you began this lesson. Type **stopwatch()** to see how long you’ve been working!

```
stopwatch()
```

Hide

END