

```

1  Course: Exploratory_Data_Analysis
2  Lesson: Clustering_Example
3
4
5  - Class: text
6  Output: "Clustering_Example. (Slides for this and other Data Science courses may be
found at github https://github.com/DataScienceSpecialization/courses/. If you care to
use them, they must be downloaded as a zip file and viewed locally. This lesson
corresponds to 04_ExploratoryAnalysis/clusteringExample.)"
7
8
9  - Class: text
10 Output: In this lesson we'll apply some of the analytic techniques we learned in
this course to data from the University of California, Irvine. Specifically, the data
we'll use is from UCI's Center for Machine Learning and Intelligent Systems. You can
find out more about the data at
http://archive.ics.uci.edu/ml/datasets/Human+Activity+Recognition+Using+Smartphones.
As this address indicates, the data involves smartphones and recognizing human
activity. Cool, right?
11
12 - Class: text
13 Output: Our goal is to show you how to use exploratory data analysis to point you in
fruitful directions of research, that is, towards answerable questions. Exploratory
data analysis is a "rough cut" or filter which helps you to find the most beneficial
areas of questioning so you can set your priorities accordingly.
14
15 - Class: text
16 Output: We also hope to show you that "real-world" research isn't always neat and
well-defined like textbook questions with clearcut answers.
17
18 - Class: cmd_question
19 Output: We've loaded data from this study for you in a matrix called ssd. Run the R
command dim now to see its dimensions.
20 CorrectAnswer: dim(ssd)
21 AnswerTests: omnittest(correctExpr='dim(ssd)')
22 Hint: Type dim(ssd) at the command prompt.
23
24 - Class: text
25 Output: Wow - ssd is pretty big, 7352 observations, each of 563 variables. Don't
worry we'll only use a small portion of this "Human Activity Recognition database".
26
27 - Class: text
28 Output: The study creating this database involved 30 volunteers "performing
activities of daily living (ADL) while carrying a waist-mounted smartphone with
embedded inertial sensors. ... Each person performed six activities ... wearing a
smartphone (Samsung Galaxy S II) on the waist. ... The experiments have been
video-recorded to label the data manually. The obtained dataset has been randomly
partitioned into two sets, where 70% of the volunteers was selected for generating
the training data and 30% the test data."
29
30 - Class: cmd_question
31 Output: Use the R command names with just the last two columns (562 and 563) of ssd
to see what data they contain.
32 CorrectAnswer: names(ssd[562:563])
33 AnswerTests:
ANY_of_exprs('names(ssd[562:563])', 'names(ssd[,562:563])', 'names(ssd[,c(562,563)])', 'na
mes(ssd[c(562,563)])', 'names(ssd[c(562:563)])', 'names(ssd[,c(562:563)])')
34 Hint: Type names(ssd[562:563]) at the command prompt.
35
36 - Class: cmd_question
37 Output: These last 2 columns contain subject and activity information. We saw above
that the gathered data had "been randomly partitioned into two sets, where 70% of
the volunteers was selected for generating the training data and 30% the test data."
Run the R command table with ssd$subject as its argument to see if the data in ssd
contains training or test data.
38 CorrectAnswer: table(ssd$subject)
39 AnswerTests: omnittest(correctExpr='table(ssd$subject)')
40 Hint: Type table(ssd$subject) at the command prompt.data."
41

```

```

42 - Class: mult_question
43 Output: From the number of subjects, would you infer that ssd contains training or
test data?
44 AnswerChoices: training; test
45 CorrectAnswer: training
46 AnswerTests: omnitest(correctVal='training')
47 Hint: Count the number of subjects represented here. Does this represent 70% or 30%
of the total subject population?
48
49 - Class: mult_question
50 Output: So ssd contains only training data. If you ran the R command sum with
table(ssd$subject) as its argument, what would the number you get back represent?
51 AnswerChoices: the number of rows in ssd; the number of columns in ssd; the number
of rows and columns of ssd; Huh?
52 CorrectAnswer: the number of rows in ssd
53 AnswerTests: omnitest(correctVal='the number of rows in ssd')
54 Hint: Each row was labeled with one subject and the output from table(ssd$subject)
told you how many rows each subject contributed to the study.
55
56 - Class: cmd_question
57 Output: Try it now (running sum on table(ssd$subject)) to see if you get 7352, the
number of rows in ssd, as a result.
58 CorrectAnswer: sum(table(ssd$subject))
59 AnswerTests: omnitest(correctExpr='sum(table(ssd$subject))')
60 Hint: Type sum(table(ssd$subject)) at the command prompt.
61
62 - Class: cmd_question
63 Output: So we're looking at training data from a machine learning repository. We can
infer that this data is supposed to train machines to recognize activity collected
from the accelerometers and gyroscopes built into the smartphones that the subjects
had strapped to their waists. Run the R command table on ssd$activity to see what
activities have been characterized by this data.
64 CorrectAnswer: table(ssd$activity)
65 AnswerTests: omnitest(correctExpr='table(ssd$activity)')
66 Hint: Type table(ssd$activity) at the command prompt.
67
68 - Class: mult_question
69 Output: We have 6 activities, 3 passive (laying, standing and sitting) and 3 active
which involve walking. If you ran the R command sum with table(ssd$activity) as its
argument, what would the number you get back represent?
70 AnswerChoices: the number of rows in ssd; the number of columns in ssd; the number
of rows and columns of ssd; Huh?
71 CorrectAnswer: the number of rows in ssd
72 AnswerTests: omnitest(correctVal='the number of rows in ssd')
73 Hint: Each row was labeled with one activity and the output from table(ssd$activity)
told you how many rows were associated with each activity in the study.
74
75 - Class: text
76 Output: Because it's training data, each row is labeled with the correct activity
(from the 6 possible) and associated with the column measurements (from the
accelerometer and gyroscope). We're interested in questions such as, "Is the
correlation between the measurements and activities good enough to train a machine?"
so that "Given a set of 561 measurements, would a trained machine be able to
determine which of the 6 activities the person was doing?"
77
78 - Class: cmd_question
79 Output: First, let's massage the data a little so it's easier to work with. We've
already run the R command transform on the data so that activities are factors. This
will let us color code them when we generate plots. Let's look at only the first
subject (numbered 1). Create the variable sub1 by assigning to it the output of the R
command subset with ssd as the first argument and the boolean, subject equal to 1,
as the second.
80 CorrectAnswer: sub1 <- subset(ssd, subject == 1)
81 AnswerTests: expr_creates_var("sub1"); omnitest(correctExpr='sub1 <- subset(ssd,
subject == 1)')
82 Hint: Type sub1 <- subset(ssd, subject == 1) at the command prompt.
83
84 - Class: cmd_question
85 Output: Look at the dimensions of sub1 now.

```

```

86 CorrectAnswer: dim(sub1)
87 AnswerTests: omnitest(correctExpr='dim(sub1)')
88 Hint: Type dim(sub1) at the command prompt.
89
90 - Class: cmd_question
91 Output: So sub1 has fewer than 400 rows now, but still a lot of columns which contain
measurements. Use names on the first 12 columns of sub1 to see what kind of data we
have.
92 CorrectAnswer: names(sub1[1:12])
93 AnswerTests:
ANY_of_exprs('names(sub1[1:12])','names(sub1[,1:12])','names(sub1)[1:12]','names(sub1[c
(1:12)])','names(sub1[,c(1:12)])')
94 Hint: Type names(sub1[1:12]) at the command prompt.
95
96 - Class: cmd_question
97 Output: We see X, Y, and Z (3 dimensions) of different aspects of body acceleration
measurements, such as mean and standard deviation. Let's do some comparisons of
activities now by looking at plots of mean body acceleration in the X and Y
directions. Call the function myedit with the string "showXY.R" to see the code
generating the plots. Make sure your cursor is back in the console window before you
hit any more buttons.
98 CorrectAnswer: myedit("showXY.R")
99 AnswerTests: omnitest(correctExpr='myedit("showXY.R")')
100 Hint: Type myedit("showXY.R") at the command prompt.
101
102 - Class: figure
103 Output: You see both the code and its output! The plots are a little squished, but
we see that the active activities related to walking (shown in the two blues and
magenta) show more variability than the passive activities (shown in black, red, and
green), particularly in the X dimension.
104 Figure: showXY.R
105 FigureType: new
106
107 - Class: cmd_question
108 Output: The colors are a little hard to distinguish. Just for fun, call the function
showMe (we used it in the Working_with_Colors lesson) which displays color vectors.
Use the vector 1:6 as its argument, and hopefully this will clarify the colors you
see in the XY comparison plot.
109 CorrectAnswer: showMe(1:6)
110 AnswerTests: ANY_of_exprs("showMe(1:6)","showMe(c(1:6))")
111 Hint: Type showMe(1:6) at the command prompt.
112
113 - Class: text
114 Output: Nice! We just wanted to show you the beauty and difference in colors. The
colors at the bottom, black, red and green, mark the passive activities, while the
true blues and magenta near the top show the walking activities. Let's try clustering
to see if we can distinguish the activities more.
115
116 - Class: cmd_question
117 Output: We'll still focus on the 3 dimensions of mean acceleration. (The plot we just
saw looked at the first 2 dimensions.) Create a distance matrix, mdist, of the first
3 columns of sub1, by using the R command dist. Use the x[,1:3] notation to specify
the columns.
118 CorrectAnswer: mdist <- dist(sub1[,1:3])
119 AnswerTests: expr_creates_var("mdist"); ANY_of_exprs('mdist <-
dist(sub1[,1:3])','mdist <- dist(sub1[,c(1:3)])')
120 Hint: Type mdist <- dist(sub1[,1:3]) the command prompt.
121
122 - Class: cmd_question
123 Output: Now create the variable hclustering by calling the R command hclust and
passing it mdist as an argument. This will use the Euclidean distance as its default
metric.
124 CorrectAnswer: hclustering <- hclust(mdist)
125 AnswerTests: expr_creates_var("hclustering"); omnitest(correctExpr='hclustering <-
hclust(mdist)')
126 Hint: Type hclustering <- hclust(mdist) the command prompt.
127
128 - Class: cmd_question
129 Output: Now call the pretty plotting function (which we've already sourced) myplclust

```

```

with 2 arguments. The first is hclustering, and the second is the argument lab.col
set equal to unclass(sub1$activity).
130 CorrectAnswer: myplclust(hclustering, lab.col = unclass(sub1$activity))
131 AnswerTests: omnitest(correctExpr='myplclust(hclustering, lab.col =
unclass(sub1$activity))')
132 Hint: Type myplclust(hclustering, lab.col = unclass(sub1$activity)) the command prompt.
133
134 - Class: text
135 Output: Well that dendrogram doesn't look too helpful, does it? There's no clear
grouping of colors, except that active colors (blues and magenta) are near each other
as are the passive (black, red, and green). So average acceleration doesn't tell us
much. How about maximum acceleration? Let's look at that for the first subject (in
our array sub1) for the X and Y dimensions. These are in column 10 and 11.
136
137 - Class: figure
138 Output: Here they are plotted side by side, X dimension on the left and Y on the
right. The x-axis of each show the 300+ observations and the y-axis indicates the
maximum acceleration.
139 Figure: showMax.R
140 FigureType: new
141
142 - Class: mult_question
143 Output: From the 2 plots, what separation, if any, do you see?
144 AnswerChoices: passive activities mostly fall below the walking activities; laying
generates the most acceleration in the X dimension; passive activities generate the
most acceleration; there is no pattern
145 CorrectAnswer: passive activities mostly fall below the walking activities
146 AnswerTests: omnitest(correctVal='passive activities mostly fall below the walking
activities')
147 Hint: There is a pattern. Which choice makes the most obvious sense?
148
149 - Class: cmd_question
150 Output: Finally we're seeing something vaguely interesting! Let's focus then on the 3
dimensions of maximum acceleration, stored in columns 10 through 12 of sub1. Create a
new distance matrix, mdist, of these 3 columns of sub1, by using the R command dist.
Again, use the x[,10:12] notation to catch the columns.
151 CorrectAnswer: mdist <- dist(sub1[,10:12])
152 AnswerTests: expr_creates_var("mdist"); ANY_of_exprs('mdist <-
dist(sub1[,10:12])', 'mdist <- dist(sub1[,c(10:12)])')
153 Hint: Type mdist <- dist(sub1[,10:12]) the command prompt.
154
155 - Class: cmd_question
156 Output: Now create the variable hclustering by calling hclust with mdist as the
argument.
157 CorrectAnswer: hclustering <- hclust(mdist)
158 AnswerTests: expr_creates_var("hclustering"); omnitest(correctExpr='hclustering <-
hclust(mdist)')
159 Hint: Type hclustering <- hclust(mdist) the command prompt.
160
161 - Class: cmd_question
162 Output: Again, call the myplclust with 2 arguments. The first is hclustering, and
the second is the argument lab.col set equal to unclass(sub1$activity).
163 CorrectAnswer: myplclust(hclustering, lab.col = unclass(sub1$activity))
164 AnswerTests: omnitest(correctExpr='myplclust(hclustering, lab.col =
unclass(sub1$activity))')
165 Hint: Type myplclust(hclustering, lab.col = unclass(sub1$activity)) the command prompt.
166
167 - Class: text
168 Output: Now we see clearly that the data splits into 2 clusters, active and passive
activities. Moreover, the light blue (walking down) is clearly distinct from the
other walking activities. The dark blue (walking level) also seems to be somewhat
clustered. The passive activities, however, seem all jumbled together with no clear
pattern visible.
169
170 - Class: cmd_question
171 Output: Let's try some SVD now. Create the variable svd1 by assigning to it the
output of a call to the R command svd. The argument to svd should be
scale(sub1[, -c(562,563)]). This will remove the last 2 columns from sub1 and scale
the data. Recall that the last 2 columns contain activity and subject information

```

```

172 which we won't need.
173 CorrectAnswer: svd1 <- svd(scale(sub1[, -c(562, 563)]))
174 AnswerTests: expr_creates_var("svd1"); ANY_of_exprs('svd1 <-
175 svd(scale(sub1[, -c(562, 563)]))', 'svd1 <- svd(scale(sub1[, -c(562:563)]))')
176 Hint: Type svd1 <- svd(scale(sub1[, -c(562, 563)])) the command prompt.
177
178 - Class: mult_question
179 Output: To see LEFT singular vectors of sub1, which component of svd1 would we examine?
180 AnswerChoices: u; v; d; x
181 CorrectAnswer: u
182 AnswerTests: omnitest(correctVal='u')
183 Hint: One of the choices isn't even part of the svd output. Recall that singular
184 value decomposition expresses the matrix X as the product of three other matrices,
185 X=UDV. Which of these is leftmost?
186
187 - Class: cmd_question
188 Output: Call the R command dim with svd1$u as an argument.
189 CorrectAnswer: dim(svd1$u)
190 AnswerTests: omnitest(correctExpr='dim(svd1$u)')
191 Hint: Type dim(svd1$u) at the command prompt.
192
193 - Class: text
194 Output: We see that the u matrix is a 347 by 347 matrix. Each row in u corresponds to
195 a row in the matrix sub1. Recall that in sub1 each row has an associated activity.
196
197 - Class: figure
198 Output: Here we're looking at the 2 left singular vectors of svd1 (the first 2
199 columns of svd1$u). Each entry of the columns belongs to a particular row with one of
200 the 6 activities assigned to it. We see the activities distinguished by color.
201 Moving from left to right, the first section of rows are green (standing), the second
202 red (sitting), the third black (laying), etc. The first column of u shows separation
203 of the nonmoving (black, red, and green) from the walking activities. The second
204 column is harder to interpret. However, the magenta cluster, which represents walking
205 up, seems separate from the others.
206 Figure: showU2.R
207 FigureType: new
208
209 - Class: text
210 Output: We'll try to figure out why that is. To do that we'll have to find which of
211 the 500+ measurements (represented by the columns of sub1) contributes to the
212 variation of that component. Since we're interested in sub1 columns, we'll look at
213 the RIGHT singular vectors (the columns of svd1$v), and in particular, the second one
214 since the separation of the magenta cluster stood out in the second column of svd1$u.
215
216 - Class: figure
217 Output: Here's a plot of the second column of svd1$v. We used transparency in our
218 plotting but nothing clearly stands out here. Let's use clustering to find the
219 feature (out of the 500+) which contributes the most to the variation of this second
220 column of svd1$v.
221 Figure: showV2.R
222 FigureType: new
223
224 - Class: cmd_question
225 Output: Create the variable maxCon by assigning to it the output of the R command
226 which.max using the second column of svd1$v as an argument.
227 CorrectAnswer: maxCon <- which.max(svd1$v[, 2])
228 AnswerTests: expr_creates_var("maxCon"); omnitest(correctExpr='maxCon <-
229 which.max(svd1$v[, 2])')
230 Hint: Type maxCon <- which.max(svd1$v[, 2]) at the command prompt.
231
232 - Class: cmd_question
233 Output: Now create a distance matrix mdist by assigning to it the output of the R
234 command dist using 4 columns of sub1 as the arguments. These 4 columns are 10
235 through 12 (10:12) and maxCon. Recall that you'll have to concatenate these 2 column
236 expressions when specifying them.
237 CorrectAnswer: mdist <- dist(sub1[, c(10:12, maxCon)])
238 AnswerTests: expr_creates_var("mdist"); omnitest(correctExpr='mdist <-
239 dist(sub1[, c(10:12, maxCon)])')
240 Hint: Type mdist <- dist(sub1[, c(10:12, maxCon)]) at the command prompt.

```



```

216
217 - Class: cmd_question
218 Output: Now create hclustering, the output of the R command hclust using mdist as the
      argument.
219 CorrectAnswer: hclustering <- hclust(mdist)
220 AnswerTests: expr_creates_var("hclustering"); omnitest(correctExpr='hclustering <-
      hclust(mdist)')
221 Hint: Type hclustering <- hclust(mdist) at the command prompt.
222
223 - Class: cmd_question
224 Output: Call the myplclust with 2 arguments, hclustering, and lab.col set equal to
      unclass(sub1$activity).
225 CorrectAnswer: myplclust(hclustering, lab.col = unclass(sub1$activity))
226 AnswerTests: omnitest(correctExpr='myplclust(hclustering, lab.col =
      unclass(sub1$activity))')
227 Hint: Type myplclust(hclustering, lab.col = unclass(sub1$activity)) at the command
      prompt.
228
229 - Class: text
230 Output: Now we see some real separation. Magenta (walking up) is on the far left, and
      the two other walking activities, the two blues, are on the far right, but in
      separate clusters from one another. The nonmoving activities still are jumbled
      together.
231
232 - Class: cmd_question
233 Output: Run the R command names with the argument sub1[maxCon] to see what
      measurement is associated with this maximum contributor.
234 CorrectAnswer: names(sub1[maxCon])
235 AnswerTests: ANY_of_exprs('names(sub1[maxCon])', 'names(sub1)[maxCon]')
236 Hint: Type names(sub1[maxCon]) or names(sub1)[maxCon] at the command prompt.
237
238 - Class: text
239 Output: So the mean body acceleration in the frequency domain in the Z direction is
      the main contributor to this clustering phenomenon we're seeing. Let's move on to
      k-means clustering to see if this technique can distinguish between the activities.
240
241 - Class: cmd_question
242 Output: Create the variable kClust by assigning to it the output of the R command
      kmeans with 2 arguments. The first is sub1 with the last 2 columns removed. (Recall
      these don't have pertinent information for clustering analysis.) The second argument
      to kmeans is centers set equal to 6, the number of activities we know we have.
243 CorrectAnswer: kClust <- kmeans(sub1[, -c(562, 563)], centers = 6)
244 AnswerTests: expr_creates_var("kClust"); ANY_of_exprs('kClust <- kmeans(sub1[,
      -c(562, 563)], centers = 6)', 'kClust <- kmeans(sub1[, -c(562:563)], centers = 6)')
245 Hint: Type kClust <- kmeans(sub1[, -c(562, 563)], centers = 6) the command prompt.
246
247 - Class: cmd_question
248 Output: Recall that without specifying coordinates for the cluster centroids (as we
      did), kmeans will generate starting points randomly. Here we did only 1 random start
      (the default). To see the output, run the R command table with 2 arguments. The first
      is kClust$cluster (part of the output from kmeans), and the second is sub1$activity.
249 CorrectAnswer: table(kClust$cluster, sub1$activity)
250 AnswerTests: omnitest(correctExpr='table(kClust$cluster, sub1$activity)')
251 Hint: Type table(kClust$cluster, sub1$activity) the command prompt.
252
253 - Class: text
254 Output: Your exact output will depend on the state of your random number generator.
      We notice that when we just run with 1 random start, the clusters tend to group the
      nonmoving activities together in one cluster. The walking activities seem to cluster
      individually by themselves. You could run the call to kmeans with one random start
      again and you'll probably get a slightly different result, but....
255
256 - Class: cmd_question
257 Output: ... instead call kmeans with 3 arguments, the last of which will tell it to
      try more random starts and return the best one. The first 2 arguments should be the
      same as before (sub1 with the last 2 columns removed and centers set equal to 6). The
      third is nstart set equal to 100. Put the result in kClust again.
258 CorrectAnswer: kClust <- kmeans(sub1[, -c(562, 563)], centers = 6, nstart=100)
259 AnswerTests: expr_creates_var("kClust"); ANY_of_exprs('kClust <- kmeans(sub1[,

```

```

- c(562, 563)], centers = 6, nstart=100)', 'kClust <- kmeans(sub1[, -c(562:563)],
centers = 6, nstart=100)')
260 Hint: Type kClust <- kmeans(sub1[, -c(562, 563)], centers = 6, nstart=100) the
command prompt.
261
262 - Class: cmd_question
263 Output: Again, run the R command table with 2 arguments. The first is kClust$cluster
(part of the output from kmeans), and the second is sub1$activity.
CorrectAnswer: table(kClust$cluster, sub1$activity)
264 AnswerTests: omnitest(correctExpr='table(kClust$cluster, sub1$activity)')
265 Hint: Type table(kClust$cluster, sub1$activity) the command prompt.
266
267 - Class: text
268 Output: We see that even with 100 random starts, the passive activities tend to
cluster together. One of the clusters contains only laying, but in another cluster,
standing and sitting group together.
269
270
271 - Class: cmd_question
272 Output: Use dim to find the dimensions of kClust's centers. Use the x$y notation to
access them.
CorrectAnswer: dim(kClust$centers)
273 AnswerTests: omnitest(correctExpr='dim(kClust$centers)')
274 Hint: Type dim(kClust$centers) the command prompt.
275
276 - Class: text
277 Output: So the centers are a 6 by 561 array. Sometimes it's a good idea to look at
the features (columns) of these centers to see if any dominate.
278
279 - Class: cmd_question
280 Output: Create the variable laying and assign to it the output of the call to the R
command which with the argument kClust$size==29.
CorrectAnswer: laying <- which(kClust$size==29)
282 AnswerTests: expr_creates_var("laying"); omnitest(correctExpr='laying <-
which(kClust$size==29)')
283 Hint: Type laying <- which(kClust$size==29) the command prompt.
284
285 - Class: cmd_question
286 Output: Now call plot with 3 arguments. The first is kClust$centers[laying,1:12],
and the second is pch set to 19. The third is ylab set equal to "Laying Cluster"
CorrectAnswer: plot(kClust$centers[laying, 1:12], pch=19, ylab="Laying Cluster")
288 AnswerTests: omnitest(correctExpr='plot(kClust$centers[laying,
1:12], pch=19, ylab="Laying Cluster")')
289 Hint: Type plot(kClust$centers[laying, 1:12], pch=19, ylab="Laying Cluster") the
command prompt.
290
291 - Class: cmd_question
292 Output: We see the first 3 columns dominate this cluster center. Run names with the
first 3 columns of sub1 as the argument to remind yourself of what these columns
contain.
293 CorrectAnswer: names(sub1[, 1:3])
294 AnswerTests:
ANY_of_exprs('names(sub1[, 1:3])', 'names(sub1[, c(1:3)])', 'names(sub1[, c(1,2,3)])', 'names
(sub1[c(1,2,3)])', 'names(sub1[1:3])', 'names(sub1[c(1:3)])', 'names(sub1)[c(1:3)]', 'names
(sub1)[c(1,2,3)]', 'names(sub1)[1:3]')
296 Hint: Type names(sub1[, 1:3]) the command prompt.
297
298 - Class: text
299 Output: So the 3 directions of mean body acceleration seem to have the biggest effect
on laying.
300
301 - Class: cmd_question
302 Output: Create the variable walkdown and assign to it the output of the call to the
R command which with the argument kClust$size==49.
CorrectAnswer: walkdown <- which(kClust$size==49)
303 AnswerTests: expr_creates_var("walkdown"); omnitest(correctExpr='walkdown <-
which(kClust$size==49)')
304 Hint: Type walkdown <- which(kClust$size==49) the command prompt.
305
306 - Class: cmd_question
307

```

```

308 Output: Now call plot with 3 arguments. The first is kClust$centers[walkdown,1:12],
and the second is pch set to 19. The third is ylab set equal to "Walkdown Cluster"
309 CorrectAnswer: plot(kClust$centers[walkdown, 1:12],pch=19,ylab="Walkdown Cluster")
310 AnswerTests: omnitest(correctExpr='plot(kClust$centers[walkdown,
1:12],pch=19,ylab="Walkdown Cluster")')
311 Hint: Type plot(kClust$centers[walkdown, 1:12],pch=19,ylab="Walkdown Cluster") the
command prompt.
312
313 - Class: text
314 Output: We see an interesting pattern here. From left to right, looking at the 12
acceleration measurements in groups of 3, the points decrease in value. The X
direction dominates, followed by Y then Z. This might tell us something more about
the walking down activity.
315
316 - Class: text
317 Output: We'll wrap up here and hope this example convinced you that real world
analysis can be frustrating sometimes and not always obvious. You might have to try
several techniques of exploratory data analysis before you hit one that pays off and
leads you to the questions that will be the most promising to explore.
318
319 - Class: text
320 Output: We saw here that the sensor measurements were pretty good at discriminating
between the 3 walking activities, but the passive activities were harder to
distinguish from one another. These might require more analysis or an entirely
different set of sensory measurements.
321
322 - Class: text
323 Output: Congratulations! We hope you enjoyed the 6 activities and 500+ features of
this lesson.
324
325 - Class: mult_question
326 Output: "Would you like to receive credit for completing this course on
327 Coursera.org?"
328 CorrectAnswer: NULL
329 AnswerChoices: Yes;No
330 AnswerTests: coursera_on_demand()
331 Hint: ""
332

```