

```

1  Course: Exploratory_Data_Analysis
2  Lesson: CaseStudy
3
4
5  - Class: text
6  Output: "CaseStudy. (Slides for this and other Data Science courses may be found at
github https://github.com/DataScienceSpecialization/courses/. If you care to use
them, they must be downloaded as a zip file and viewed locally. This lesson
corresponds to 04_ExploratoryAnalysis/CaseStudy.)"
7
8  - Class: text
9  Output: In this lesson we'll apply some of the techniques we learned in this course
to study air pollution data, specifically particulate matter (we'll call it pm25
sometimes), collected by the U.S. Environmental Protection Agency. This website
https://www.health.ny.gov/environmental/indoors/air/pmq\_a.htm from New York State
offers some basic information on this topic if you're interested.
10
11 - Class: text
12 Output: Particulate matter (less than 2.5 microns in diameter) is a fancy name for
dust, and breathing in dust might pose health hazards to the population. We'll study
data from two years, 1999 (when monitoring of particulate matter started) and 2012.
Our goal is to see if there's been a noticeable decline in this type of air pollution
between these two years.
13
14 - Class: cmd_question
15 Output: We've read in 2 large zipped files for you using the R command read.table
(which is smart enough to unzip the files). We stored the 1999 data in the array pm0
for you. Run the R command dim now to see its dimensions.
16 CorrectAnswer: dim(pm0)
17 AnswerTests: omnitest(correctExpr='dim(pm0)')
18 Hint: Type dim(pm0) at the command prompt.
19
20 - Class: cmd_question
21 Output: We see that pm0 has over 117000 lines, each containing 5 columns. In the
original file, at the EPA website, each row had 28 columns, but since we'll be using
only a few of these, we've created and read in a somewhat smaller file. Run head on
pm0 now to see what the first few lines look like.
22 CorrectAnswer: head(pm0)
23 AnswerTests: omnitest(correctExpr='head(pm0)')
24 Hint: Type head(pm0) at the command prompt.
25
26 - Class: text
27 Output: We see there's some missing data, but we won't worry about that now. We also
see that the column names, V1, V2, etc., are not informative. However, we know that
the first line of the original file (a comment) explained what information the
columns contained.
28
29 - Class: cmd_question
30 Output: We created the variable cnames containing the 28 column names of the
original file. Take a look at the column names now.
31 CorrectAnswer: cnames
32 AnswerTests: ANY_of_exprs('cnames', 'print(cnames)')
33 Hint: Type cnames or print(cnames) at the command prompt.
34
35 - Class: cmd_question
36 Output: We see that the 28 column names look all jumbled together even though
they're separated by "|" characters, so let's fix this. Reassign to cnames the output
of a call to strsplit (string split) with 3 arguments. The first is cnames, the pipe
symbol '|' is the second (use the quotation marks), and the third is the argument
fixed set to TRUE. Try this now.
37 CorrectAnswer: cnames <- strsplit(cnames, "|", fixed = TRUE)
38 AnswerTests: any_of_exprs('cnames <- strsplit(cnames, "|", fixed = T)', 'cnames <-
strsplit(cnames, "|", fixed = TRUE)')
39 Hint: Type cnames <- strsplit(cnames, "|", fixed = TRUE) at the command prompt.
40
41 - Class: cmd_question
42 Output: The variable cnames now holds a list of the column headings. Take another
look at the column names.
43 CorrectAnswer: cnames

```

```

44 AnswerTests: ANY_of_exprs('cnames','print(cnames)')
45 Hint: Type cnames or print(cnames) at the command prompt.
46
47 - Class: cmd_question
48 Output: Nice, but we don't need all these. Assign to names(pm0) the output of a
call to the function make.names with cnames[[1]][wcol] as the argument. The variable
wcol holds the indices of the 5 columns we selected (from the 28) to use in this
lesson, so those are the column names we'll need. As the name suggests, the function
"makes syntactically valid names".
49 CorrectAnswer: names(pm0) <- make.names(cnames[[1]][wcol])
50 AnswerTests: omnitest(correctExpr='names(pm0) <- make.names(cnames[[1]][wcol])')
51 Hint: Type names(pm0) <- make.names(cnames[[1]][wcol]) at the command prompt.
52
53 - Class: cmd_question
54 Output: Now re-run head on pm0 now to see if the column names have been put in place.
55 CorrectAnswer: head(pm0)
56 AnswerTests: omnitest(correctExpr='head(pm0)')
57 Hint: Type head(pm0) at the command prompt.
58
59 - Class: cmd_question
60 Output: Now it's clearer what information each column of pm0 holds. The measurements
of particulate matter (pm25) are in the column named Sample.Value. Assign this
component of pm0 to the variable x0. Use the m$n notation.
61 CorrectAnswer: x0 <- pm0$Sample.Value
62 AnswerTests: expr_creates_var("x0"); omnitest(correctExpr='x0 <- pm0$Sample.Value')
63 Hint: Type x0 <- pm0$Sample.Value at the command prompt.
64
65 - Class: cmd_question
66 Output: Call the R command str with x0 as its argument to see x0's structure.
67 CorrectAnswer: str(x0)
68 AnswerTests: omnitest(correctExpr='str(x0)')
69 Hint: Type str(x0) at the command prompt.
70
71 - Class: cmd_question
72 Output: We see that x0 is a numeric vector (of length 117000+) with at least the
first 3 values missing. Exactly what percentage of values are missing in this
vector? Use the R function mean with is.na(x0) as an argument to see what percentage
of values are missing (NA) in x0.
73 CorrectAnswer: mean(is.na(x0))
74 AnswerTests: omnitest(correctExpr='mean(is.na(x0))')
75 Hint: Type mean(is.na(x0)) at the command prompt.
76
77 - Class: text
78 Output: So a little over 11% of the 117000+ are missing. We'll keep that in mind. Now
let's start processing the 2012 data which we stored for you in the array pm1.
79
80 - Class: cmd_question
81 Output: We'll repeat what we did for pm0, except a little more efficiently. First
assign the output of make.names(cnames[[1]][wcol]) to names(pm1).
82 CorrectAnswer: names(pm1) <- make.names(cnames[[1]][wcol])
83 AnswerTests: omnitest(correctExpr='names(pm1) <- make.names(cnames[[1]][wcol])')
84 Hint: Type names(pm1) <- make.names(cnames[[1]][wcol]) at the command prompt.
85
86 - Class: cmd_question
87 Output: Find the dimensions of pm1 with the command dim.
88 CorrectAnswer: dim(pm1)
89 AnswerTests: omnitest(correctExpr='dim(pm1)')
90 Hint: Type dim(pm1) at the command prompt.
91
92 - Class: text
93 Output: Wow! Over 1.3 million entries. Particulate matter was first collected in 1999
so perhaps there weren't as many sensors collecting data then as in 2012 when the
program was more mature. If you ran head on pm1 you'd see that it looks just like
pm0. We'll move on though.
94
95 - Class: cmd_question
96 Output: Create the variable x1 by assigning to it the Sample.Value component of pm1.
97 CorrectAnswer: x1 <- pm1$Sample.Value
98 AnswerTests: expr_creates_var("x1"); omnitest(correctExpr='x1 <- pm1$Sample.Value')

```

```

99     Hint: Type x1 <- pml$Sample.Value at the command prompt.
100
101 - Class: cmd_question
102 Output: Now let's see what percentage of values are missing in x1. As before, use
the R function mean with is.na(x1) as an argument to find out.
103 CorrectAnswer: mean(is.na(x1))
104 AnswerTests: omnitest(correctExpr='mean(is.na(x1))')
105 Hint: Type mean(is.na(x1)) at the command prompt.
106
107 - Class: text
108 Output: So only 5.6% of the particulate matter measurements are missing. That's about
half the percentage as in 1999.
109
110 - Class: cmd_question
111 Output: Now let's look at summaries (using the summary command) for both datasets.
First, x0.
112 CorrectAnswer: summary(x0)
113 AnswerTests: omnitest(correctExpr='summary(x0)')
114 Hint: Type summary(x0) at the command prompt.
115
116 - Class: cmd_question
117 Output: The numbers in the vectors x0 and x1 represent measurements taken in
micrograms per cubic meter. Now look at the summary of x1.
118 CorrectAnswer: summary(x1)
119 AnswerTests: omnitest(correctExpr='summary(x1)')
120 Hint: Type summary(x1) at the command prompt.
121
122 - Class: text
123 Output: We see that both the median and the mean of measured particulate matter have
declined from 1999 to 2012. In fact, all of the measurements, except for the maximum
and missing values (Max and NA's), have decreased. Even the Min has gone down from 0
to -10.00! We'll address what a negative measurement might mean a little later. Note
that the Max has increased from 157 in 1999 to 909 in 2012. This is quite high and
might reflect an error in the table or malfunctions in some monitors.
124
125 - Class: cmd_question
126 Output: Call the boxplot function with 2 arguments, x0 and x1.
127 CorrectAnswer: boxplot(x0, x1)
128 AnswerTests: omnitest(correctExpr='boxplot(x0, x1)')
129 Hint: Type boxplot(x0, x1) at the command prompt.
130
131 - Class: cmd_question
132 Output: Huh? Did somebody step on the boxes? It's hard to see what's going on here.
There are so many values outside the boxes and the range of x1 is so big that the
boxes are flattened. It might be more informative to call boxplot on the logs (base
10) of x0 and x1. Do this now using log10(x0) and log10(x1) as the 2 arguments.
133 CorrectAnswer: boxplot(log10(x0), log10(x1))
134 AnswerTests: omnitest(correctExpr='boxplot(log10(x0), log10(x1))')
135 Hint: Type boxplot(log10(x0), log10(x1)) at the command prompt.
136
137 - Class: text
138 Output: A bonus! Not only do we get a better looking boxplot we also get some
warnings from R in Red. These let us know that some values in x0 and x1 were
"unloggable", no doubt the 0 (Min) we saw in the summary of x0 and the negative
values we saw in the Min of the summary of x1.
139
140 - Class: mult_question
141 Output: From the boxplot (x0 on the left and x1 on the right), what can you say about
the data?
142 AnswerChoices: The median of x1 is less than the median of x0; The mean of x1 is
less than the mean of x0; The range of x0 is greater than the range of x1; The boxes
are too small to interpret
143 CorrectAnswer: The median of x1 is less than the median of x0
144 AnswerTests: omnitest(correctVal='The median of x1 is less than the median of x0')
145 Hint: Recall that the horizontal lines inside the boxes represent the medians of the
two datasets.
146
147 - Class: text
148 Output: Let's return to the question of the negative values in x1. Let's count how

```

many negative values there are. We'll do this in a few steps.

```
149
150 - Class: cmd_question
151 Output: First, form the vector negative by assigning to it the boolean x1<0.
152 CorrectAnswer: negative <- x1 < 0
153 AnswerTests: expr_creates_var("negative"); omnitest(correctExpr='negative <- x1 < 0')
154 Hint: Type negative <- x1 < 0 at the command prompt.
155
156 - Class: cmd_question
157 Output: Now run the R command sum with 2 arguments. The first is negative, and the
    second is na.rm set equal to TRUE. This tells sum to ignore the missing values in
    negative.
158 CorrectAnswer: sum(negative, na.rm = TRUE)
159 AnswerTests: omnitest(correctExpr='sum(negative, na.rm = TRUE)')
160 Hint: Type sum(negative, na.rm = TRUE) at the command prompt.
161
162 - Class: cmd_question
163 Output: So there are over 26000 negative values. Sounds like a lot. Is it? Run the R
    command mean with same 2 arguments you just used with the call to sum. This will tell
    us a percentage.
164 CorrectAnswer: mean(negative, na.rm = TRUE)
165 AnswerTests: omnitest(correctExpr='mean(negative, na.rm = TRUE)')
166 Hint: Type mean(negative, na.rm = TRUE) at the command prompt.
167
168 - Class: text
169 Output: We see that just 2% of the x1 values are negative. Perhaps that's a small
    enough percentage that we can ignore them. Before we ignore them, though, let's see
    if they occur during certain times of the year.
170
171 - Class: cmd_question
172 Output: First create the array dates by assigning to it the Date component of pml.
    Remember to use the x$y notation.
173 CorrectAnswer: dates <- pml$Date
174 AnswerTests: expr_creates_var("dates"); omnitest(correctExpr='dates <- pml$Date')
175 Hint: Type dates <- pml$Date at the command prompt.
176
177 - Class: cmd_question
178 Output: To see what dates looks like run the R command str on it.
179 CorrectAnswer: str(dates)
180 AnswerTests: omnitest(correctExpr='str(dates)')
181 Hint: Type str(dates) at the command prompt.
182
183 - Class: cmd_question
184 Output: We see dates is a very long vector of integers. However, the format of the
    entries is hard to read. There's no separation between the year, month, and day.
    Reassign to dates the output of a call to as.Date with the 2 arguments
    as.character(dates) as the first argument and the string "%Y%m%d" as the second.
185 CorrectAnswer: dates <- as.Date(as.character(dates), "%Y%m%d")
186 AnswerTests: expr_creates_var("dates"); omnitest(correctExpr='dates <-
    as.Date(as.character(dates), "%Y%m%d")')
187 Hint: Type dates <- as.Date(as.character(dates), "%Y%m%d") at the command prompt.
188
189 - Class: cmd_question
190 Output: Now when you run head on dates you'll see the dates in a nicer format. Try
    this now.
191 CorrectAnswer: head(dates)
192 AnswerTests: omnitest(correctExpr='head(dates)')
193 Hint: Type head(dates) at the command prompt.
194
195 - Class: cmd_question
196 Output: Let's plot a histogram of the months when the particulate matter measurements
    are negative. Run hist with 2 arguments. The first is dates[negative] and the second
    is the string "month".
197 CorrectAnswer: hist(dates[negative], "month")
198 AnswerTests: omnitest(correctExpr='hist(dates[negative], "month")')
199 Hint: Type hist(dates[negative], "month") at the command prompt.
200
201 - Class: text
202 Output: We see the bulk of the negative measurements were taken in the winter months,
```

with a spike in May. Not many of these negative measurements occurred in summer months. We can take a guess that because particulate measures tend to be low in winter and high in summer, coupled with the fact that higher densities are easier to measure, that measurement errors occurred when the values were low. For now we'll attribute these negative measurements to errors. Also, since they account for only 2% of the 2012 data, we'll ignore them.

```
203
204
205 - Class: figure
206 Output: Now we'll change focus a bit and instead of looking at all the monitors
    throughout the country and the data they recorded, we'll try to find one monitor that
    was taking measurements in both 1999 and 2012. This will allow us to control for
    different geographical and environmental variables that might have affected air
    quality in different areas. We'll narrow our search and look just at monitors in New
    York State.
207 Figure: runSites.R
208 FigureType: new
209
210 - Class: cmd_question
211 Output: We subsetting off the New York State monitor identification data for 1999 and
    2012 into 2 vectors, site0 and site1. Look at the structure of site0 now with the R
    command str.
212 CorrectAnswer: str(site0)
213 AnswerTests: obliterate("dates"); omnitest(correctExpr='str(site0)')
214 Hint: Type str(site0) at the command prompt.
215
216 - Class: text
217 Output: We see that site0 (the IDs of monitors in New York State in 1999) is a
    vector of 33 strings, each of which has the form "x.y". We've created these from the
    county codes (the x portion of the string) and the monitor IDs (the y portion). If
    you ran str on site1 you'd see 18 similar values.
218
219 - Class: cmd_question
220 Output: Use the intersect command with site0 and site1 as arguments and put the
    result in the variable both.
221 CorrectAnswer: both <- intersect(site0, site1)
222 AnswerTests: obliterate("x0"); expr_creates_var("both");
    omnitest(correctExpr='both <- intersect(site0, site1)')
223 Hint: Type both <- intersect(site0, site1) at the command prompt.
224
225 - Class: cmd_question
226 Output: Take a look at both now.
227 CorrectAnswer: both
228 AnswerTests: obliterate("x1"); ANY_of_exprs('both', 'print(both)')
229 Hint: Type both or print(both) at the command prompt.
230
231 - Class: figure
232 Output: We see that 10 monitors in New York State were active in both 1999 and 2012.
233 Figure: runModpms.R
234 FigureType: new
235
236 - Class: cmd_question
237 Output: To save you some time and typing, we modified the data frames pm0 and pm1
    slightly by adding to each of them a new component, county.site. This is just a
    concatenation of two original components County.Code and Site.ID. We did this to
    facilitate the next step which is to find out how many measurements were taken by the
    10 New York monitors working in both of the years of interest. Run head on pm0 to see
    the first few entries now.
238 CorrectAnswer: head(pm0)
239 AnswerTests: obliterate("negative"); omnitest(correctExpr='head(pm0)')
240 Hint: Type head(pm0) at the command prompt.
241
242 - Class: text
243 Output: Now pm0 and pm1 have 6 columns instead of 5, and the last column is a
    concatenation of two other columns, County and Site.
244
245 - Class: text
246 Output: Now let's see how many measurements each of the 10 New York monitors that
    were active in both 1999 and 2012 took in those years. We'll create 2 subsets (one
```

```

247     for each year), one of pm0 and the otherof pm1.
248 - Class: text
249 Output: The subsets will filter for 2 characteristics. The first is State.Code equal
to 36 (the code for New York), and the second is that the county.site (the component
we added) is in the vector both.
250
251 - Class: cmd_question
252 Output: First create the variable cnt0 by assigning to it the output of the R
command subset, called with 2 arguments. The first is pm0, and the second is a
boolean with the 2 conditions we just mentioned. Recall that the testing for equality
in a boolean requires ==, intersection of 2 boolean conditions is denoted by & and
membership by %in%.
253 CorrectAnswer: cnt0 <- subset(pm0, State.Code == 36 & county.site %in% both)
254 AnswerTests: expr_creates_var("cnt0"); omnitest(correctExpr='cnt0 <- subset(pm0,
State.Code == 36 & county.site %in% both)')
255 Hint: Type cnt0 <- subset(pm0, State.Code == 36 & county.site %in% both) at the
command prompt.
256
257 - Class: cmd_question
258 Output: Recall the last command with the up arrow, and create cnt1 (instead of
cnt0). Remember to change pm0 to pm1. Everything else can stay the same.
259 CorrectAnswer: cnt1 <- subset(pm1, State.Code == 36 & county.site %in% both)
260 AnswerTests: expr_creates_var("cnt1"); omnitest(correctExpr='cnt1 <- subset(pm1,
State.Code == 36 & county.site %in% both)')
261 Hint: Type cnt1 <- subset(pm1, State.Code == 36 & county.site %in% both) at the
command prompt.
262
263 - Class: cmd_question
264 Output: Now run the command sapply(split(cnt0, cnt0$county.site), nrow). This will
split cnt0 into several data frames according to county.site (that is, monitor IDs)
and tell us how many measurements each monitor recorded.
265 CorrectAnswer: sapply(split(cnt0, cnt0$county.site), nrow)
266 AnswerTests: omnitest(correctExpr='sapply(split(cnt0, cnt0$county.site), nrow)')
267 Hint: Type sapply(split(cnt0, cnt0$county.site), nrow) at the command prompt.
268
269 - Class: cmd_question
270 Output: Do the same for cnt1. (Recall your last command and change 2 occurrences of
cnt0 to cnt1.)
271 CorrectAnswer: sapply(split(cnt1, cnt1$county.site), nrow)
272 AnswerTests: omnitest(correctExpr='sapply(split(cnt1, cnt1$county.site), nrow)')
273 Hint: Type sapply(split(cnt1, cnt1$county.site), nrow) at the command prompt.
274
275 - Class: mult_question
276 Output: From the output of the 2 calls to sapply, which monitor is the only one whose
number of measurements increased from 1999 to 2012?
277 AnswerChoices: 29.5; 85.55; 63.2008; 101.3
278 CorrectAnswer: 85.55
279 AnswerTests: omnitest(correctVal='85.55')
280 Hint: Compare the 2 rows of numbers and look for the one with the bigger bottom number.
281
282 - Class: cmd_question
283 Output: We want to examine a monitor with a reasonable number of measurements so
let's look at the monitor with ID 63.2008. Create a variable pm0sub which is the
subset of cnt0 (this contains just New York data) which has County.Code equal to 63
and Site.ID 2008.
284 CorrectAnswer: pm0sub <- subset(cnt0, County.Code==63 & Site.ID==2008)
285 AnswerTests: expr_creates_var("pm0sub"); omnitest(correctExpr='pm0sub <-
subset(cnt0, County.Code==63 & Site.ID==2008)')
286 Hint: Type pm0sub <- subset(cnt0, County.Code==63 & Site.ID==2008) at the command
prompt.
287
288 - Class: cmd_question
289 Output: Now do the same for cnt1. Name this new variable pm1sub.
290 CorrectAnswer: pm1sub <- subset(cnt1, County.Code==63 & Site.ID==2008)
291 AnswerTests: expr_creates_var("pm1sub"); omnitest(correctExpr='pm1sub <-
subset(cnt1, County.Code==63 & Site.ID==2008)')
292 Hint: Type pm1sub <- subset(cnt1, County.Code==63 & Site.ID==2008) at the command
prompt.

```



```

293
294 - Class: mult_question
295 Output: From the output of the 2 calls to supply, how many rows will pm0sub have?
296 AnswerChoices: 122; 30; 29; 183
297 CorrectAnswer: 122
298 AnswerTests: omnitest(correctVal='122')
299 Hint: Recall monitor 63.2008 had 122 measurements in 1999 and 30 in 2012.
300
301 - Class: cmd_question
302 Output: Now we'd like to compare the pm25 measurements of this particular monitor
(63.2008) for the 2 years. First, create the vector x0sub by assigning to it the
Sample.Value component of pm0sub.
303 CorrectAnswer: x0sub <- pm0sub$Sample.Value
304 AnswerTests: expr_creates_var("x0sub"); omnitest(correctExpr='x0sub <-
pm0sub$Sample.Value')
305 Hint: Type x0sub <- pm0sub$Sample.Value at the command prompt.
306
307 - Class: cmd_question
308 Output: Similarly, create x1sub from pm1sub.
309 CorrectAnswer: x1sub <- pm1sub$Sample.Value
310 AnswerTests: expr_creates_var("x1sub"); omnitest(correctExpr='x1sub <-
pm1sub$Sample.Value')
311 Hint: Type x1sub <- pm1sub$Sample.Value at the command prompt.
312
313 - Class: cmd_question
314 Output: We'd like to make our comparison visually so we'll have to create a time
series of these pm25 measurements. First, create a dates0 variable by assigning to it
the output of a call to as.Date. This will take 2 arguments. The first is a call to
as.character with pm0sub$Date as the argument. The second is the format string
"%Y%m%d".
315 CorrectAnswer: dates0 <- as.Date(as.character(pm0sub$Date), "%Y%m%d")
316 AnswerTests: expr_creates_var("dates0"); omnitest(correctExpr='dates0 <-
as.Date(as.character(pm0sub$Date), "%Y%m%d")')
317 Hint: Type dates0 <- as.Date(as.character(pm0sub$Date), "%Y%m%d") at the command prompt.
318
319 - Class: cmd_question
320 Output: Do the same for the 2012 data. Specifically, create dates1 using pm1sub$Date
as your input.
321 CorrectAnswer: dates1 <- as.Date(as.character(pm1sub$Date), "%Y%m%d")
322 AnswerTests: expr_creates_var("dates1"); omnitest(correctExpr='dates1 <-
as.Date(as.character(pm1sub$Date), "%Y%m%d")')
323 Hint: Type dates1 <- as.Date(as.character(pm1sub$Date), "%Y%m%d") at the command prompt.
324
325 - Class: cmd_question
326 Output: Now we'll plot these 2 time series in the same panel using the base plotting
system. Call par with 2 arguments. The first is mfrow set equal to c(1,2). This will
tell the system we're plotting 2 graphs in 1 row and 2 columns. The second argument
will adjust the panel's margins. It is mar set to c(4,4,2,1).
327 CorrectAnswer: par(mfrow = c(1, 2), mar = c(4, 4, 2, 1))
328 AnswerTests: omnitest(correctExpr='par(mfrow = c(1, 2), mar = c(4, 4, 2, 1))')
329 Hint: Type par(mfrow = c(1, 2), mar = c(4, 4, 2, 1)) at the command prompt.
330
331 - Class: cmd_question
332 Output: Call plot with the 3 arguments dates0, x0sub, and pch set to 20. The first
two arguments are the x and y coordinates. This will show the pm25 values as
functions of time.
333 CorrectAnswer: plot(dates0, x0sub, pch = 20)
334 AnswerTests: omnitest(correctExpr='plot(dates0, x0sub, pch = 20)')
335 Hint: Type plot(dates0, x0sub, pch = 20) at the command prompt.
336
337 - Class: figure
338 Output: Now we'll mark the median.
339 Figure: goodPlot1.R
340
341 - Class: cmd_question
342 Output: Use abline to add a horizontal line at the median of the pm25 values. Make
the line width 2 (lwd is the argument), and when you call median with x0sub, specify
the argument na.rm to be TRUE.
343 CorrectAnswer: abline(h = median(x0sub, na.rm = TRUE), lwd=2)

```

```

344 AnswerTests: omnitest(correctExpr='abline(h = median(x0sub, na.rm = TRUE),lwd=2)')
345 Hint: Type abline(h = median(x0sub, na.rm = TRUE),lwd=2) at the command prompt.
346
347 - Class: cmd_question
348 Output: Now we'll do the same for the 2012 data. Call plot with the 3 arguments
      dates1, x1sub, and pch set to 20.
349 CorrectAnswer: plot(dates1, x1sub, pch = 20)
350 AnswerTests: omnitest(correctExpr='plot(dates1, x1sub, pch = 20)')
351 Hint: Type plot(dates1, x1sub, pch = 20) at the command prompt.
352
353 - Class: figure
354 Output: As before, we'll mark the median of this 2012 data.
355 Figure: goodPlot2.R
356
357 - Class: cmd_question
358 Output: Use abline to add a horizontal line at the median of the pm25 values. Make
      the line width 2 (lwd is the argument). Remember to specify the argument na.rm to be
      TRUE when you call median on x1sub.
359 CorrectAnswer: abline(h = median(x1sub, na.rm = TRUE),lwd=2)
360 AnswerTests: omnitest(correctExpr='abline(h = median(x1sub, na.rm = TRUE),lwd=2)')
361 Hint: Type abline(h = median(x1sub, na.rm = TRUE),lwd=2) at the command prompt.
362
363 - Class: mult_question
364 Output: Which median is larger - the one for 1999 or the one for 2012?
365 AnswerChoices: 1999; 2012
366 CorrectAnswer: 1999
367 AnswerTests: omnitest(correctVal='1999')
368 Hint: Look carefully at the numbers on the y-axis.
369
370 - Class: mult_question
371 Output: The picture makes it look like the median is higher for 2012 than 1999.
      Closer inspection shows that this isn't true. The median for 1999 is a little over 10
      micrograms per cubic meter and for 2012 its a little over 8. The plots appear this
      way because the 1999 plot ....
372 AnswerChoices: shows a bigger range of y values than the 2012 plot; displays more
      points than the 2012 plot; shows different months than those in the 2012 plot
373 CorrectAnswer: shows a bigger range of y values than the 2012 plot
374 AnswerTests: omnitest(correctVal='shows a bigger range of y values than the 2012 plot')
375 Hint: Look at the range of y values in the two plots.
376
377 - Class: cmd_question
378 Output: The 1999 plot shows a much bigger range of pm25 values on the y axis, from
      below 10 to 40, while the 2012 pm25 values are much more restricted, from around 1
      to 14. We should really plot the points of both datasets on the same range of values
      on the y axis. Create the variable rng by assigning to it the output of a call to the
      R command range with 3 arguments, x0sub, x1sub, and the boolean na.rm set to TRUE.
379 CorrectAnswer: rng <- range(x0sub,x1sub,na.rm=TRUE)
380 AnswerTests: expr_creates_var("rng"); omnitest(correctExpr='rng <-
      range(x0sub,x1sub,na.rm=TRUE)')
381 Hint: Type rng <- range(x0sub,x1sub,na.rm=TRUE) at the command prompt.
382
383 - Class: cmd_question
384 Output: Look at rng to see the values it spans.
385 CorrectAnswer: rng
386 AnswerTests: ANY_of_exprs("rng","print(rng)")
387 Hint: Type rng or print(rng) at the command prompt.
388
389 - Class: figure
390 Output: Here a new figure we've created showing the two plots side by side with the
      same range of values on the y axis. We used the argument ylim set equal to rng in our
      2 calls to plot. The improvement in the medians between 1999 and 2012 is now clear.
      Also notice that in 2012 there are no big values (above 15). This shows that not only
      is there a chronic improvement in air quality, but also there are fewer days with
      severe pollution.
391 Figure: showBoth.R
392 FigureType: new
393
394 - Class: text
395 Output: The last avenue of this data we'll explore (and we'll do it quickly) concerns

```


a comparison of all the states' mean pollution levels. This is important because the states are responsible for implementing the regulations set at the federal level by the EPA.

```
396
397 - Class: text
398 Output: Let's first gather the mean (average measurement) for each state in 1999.
Recall that the original data for this year was stored in pm0.
399
400 - Class: cmd_question
401 Output: Create the vector mn0 with a call to the R command with using 2 arguments.
The first is pm0. This is the data in which the second argument, an expression, will
be evaluated. The second argument is a call to the function tapply. This call
requires 4 arguments. Sample.Value and State.Code are the first two. We want to apply
the function mean to Sample.Value, so mean is the third argument. The fourth is
simply the boolean na.rm set to TRUE.
402 CorrectAnswer: mn0 <- with(pm0,tapply(Sample.Value,State.Code,mean,na.rm=TRUE))
403 AnswerTests: expr_creates_var("mn0"); omnitest(correctExpr='mn0 <-
with(pm0,tapply(Sample.Value,State.Code,mean,na.rm=TRUE))')
404 Hint: Type mn0 <- with(pm0,tapply(Sample.Value,State.Code,mean,na.rm=TRUE)) at the
command prompt.
405
406 - Class: cmd_question
407 Output: Call the function str with mn0 as its argument to see what it looks like.
408 CorrectAnswer: str(mn0)
409 AnswerTests: omnitest(correctExpr='str(mn0)')
410 Hint: Type str(mn0) at the command prompt.
411
412 - Class: text
413 Output: We see mn0 is a 53 long numerical vector. Why 53 if there are only 50 states?
As it happens, pm25 measurements for the District of Columbia (Washington D.C), the
Virgin Islands, and Puerto Rico are included in this data. They are coded as 11, 72,
and 78 respectively.
414
415 - Class: cmd_question
416 Output: Recall your command creating mn0 and change it to create mn1 using pm1 as the
first input to the call to with.
417 CorrectAnswer: mn1 <- with(pm1,tapply(Sample.Value,State.Code,mean,na.rm=TRUE))
418 AnswerTests: expr_creates_var("mn1"); omnitest(correctExpr='mn1 <-
with(pm1,tapply(Sample.Value,State.Code,mean,na.rm=TRUE))')
419 Hint: Type mn1 <- with(pm1,tapply(Sample.Value,State.Code,mean,na.rm=TRUE)) at the
command prompt.
420
421 - Class: cmd_question
422 Output: For fun, call the function str with mn1 as its argument.
423 CorrectAnswer: str(mn1)
424 AnswerTests: omnitest(correctExpr='str(mn1)')
425 Hint: Type str(mn1) at the command prompt.
426
427 - Class: cmd_question
428 Output: So mn1 has only 52 entries, rather than 53. We checked. There are no entries
for the Virgin Islands in 2012. Call summary now with mn0 as its input.
429 CorrectAnswer: summary(mn0)
430 AnswerTests: omnitest(correctExpr='summary(mn0)')
431 Hint: Type summary(mn0) at the command prompt.
432
433 - Class: cmd_question
434 Output: Now call summary with mn1 as its input so we can compare the two years.
435 CorrectAnswer: summary(mn1)
436 AnswerTests: omnitest(correctExpr='summary(mn1)')
437 Hint: Type summary(mn1) at the command prompt.
438
439 - Class: cmd_question
440 Output: We see that in all 6 entries, the 2012 numbers are less than those in 1999.
Now we'll create 2 new dataframes containing just the state names and their mean
measurements for each year. First, we'll do this for 1999. Create the data frame d0
by calling the function data.frame with 2 arguments. The first is state set equal to
names(mn0), and the second is mean set equal to mn0.
441 CorrectAnswer: d0 <- data.frame(state = names(mn0), mean = mn0)
442 AnswerTests: expr_creates_var("d0"); omnitest(correctExpr='d0 <- data.frame(state
```

```

= names(mn0), mean = mn0)')
443 Hint: Type d0 <- data.frame(state = names(mn0), mean = mn0) at the command prompt.
444
445 - Class: cmd_question
446 Output: Recall the last command and create d1 instead of d0 using the 2012 data.
(There'll be 3 changes of 0 to 1.)
447 CorrectAnswer: d1 <- data.frame(state = names(mn1), mean = mn1)
448 AnswerTests: expr_creates_var("d1"); omnitest(correctExpr='d1 <- data.frame(state
= names(mn1), mean = mn1)')
449 Hint: Type d1 <- data.frame(state = names(mn1), mean = mn1) at the command prompt.
450
451 - Class: cmd_question
452 Output: Create the array mrg by calling the R command merge with 3 arguments, d0, d1,
and the argument by set equal to the string "state".
453 CorrectAnswer: mrg <- merge(d0, d1, by = "state")
454 AnswerTests: expr_creates_var("mrg"); omnitest(correctExpr='mrg <- merge(d0, d1,
by = "state")')
455 Hint: Type mrg <- merge(d0, d1, by = "state") at the command prompt.
456
457 - Class: cmd_question
458 Output: Run dim with mrg as its argument to see how big it is.
459 CorrectAnswer: dim(mrg)
460 AnswerTests: omnitest(correctExpr='dim(mrg)')
461 Hint: Type dim(mrg) at the command prompt.
462
463 - Class: cmd_question
464 Output: We see merge has 52 rows and 3 columns. Since the Virgin Island data was
missing from d1, it is excluded from mrg. Look at the first few entries of mrg using
the head command.
465 CorrectAnswer: head(mrg)
466 AnswerTests: omnitest(correctExpr='head(mrg)')
467 Hint: Type head(mrg) at the command prompt.
468
469 - Class: text
470 Output: Each row of mrg has 3 entries - a state identified by number, a state mean
for 1999 (mean.x), and a state mean for 2012 (mean.y).
471
472 - Class: text
473 Output: Now we'll plot the data to see how the state means changed between the 2
years. First we'll plot the 1999 data in a single column at x=1. The y values for the
points will be the state means. Again, we'll use the R command with so we don't have
to keep typing mrg as the data environment in which to evaluate the second argument,
the call to plot. We've already reset the graphical parameters for you.
474
475 - Class: cmd_question
476 Output: For the first column of points, call with with 2 arguments. The first is mrg,
and the second is the call to plot with 3 arguments. The first of these is rep(1,52).
This tells the plot routine that the x coordinates for all 52 points are 1. The
second argument is the second column of mrg or mrg[,2] which holds the 1999 data. The
third argument is the range of x values we want, namely xlim set to c(.5,2.5). This
works since we'll be plotting 2 columns of points, one at x=1 and the other at x=2.
477 CorrectAnswer: with(mrg, plot(rep(1, 52), mrg[, 2], xlim = c(.5, 2.5)))
478 AnswerTests: omnitest(correctExpr='with(mrg, plot(rep(1, 52), mrg[, 2], xlim =
c(.5, 2.5)))')
479 Hint: Type with(mrg, plot(rep(1, 52), mrg[, 2], xlim = c(.5, 2.5))) at the command
prompt.
480
481 - Class: cmd_question
482 Output: We see a column of points at x=1 which represent the 1999 state means. For
the second column of points, again call with with 2 arguments. As before, the first
is mrg. The second, however, is a call to the function points with 2 arguments. We
need to do this since we're adding points to an already existing plot. The first
argument to points is the set of x values, rep(2,52). The second argument is the set
of y values, mrg[,3]. Of course, this is the third column of mrg. (We don't need to
specify the range of x values again.)
483 CorrectAnswer: with(mrg, points(rep(2, 52), mrg[, 3]))
484 AnswerTests: omnitest(correctExpr='with(mrg, points(rep(2, 52), mrg[, 3]))')
485 Hint: Type with(mrg, points(rep(2, 52), mrg[, 3])) at the command prompt.
486

```

```

487 - Class: cmd_question
488 Output: We see a shorter column of points at x=2. Now let's connect the dots. Use the
R function segments with 4 arguments. The first 2 are the x and y coordinates of the
1999 points and the last 2 are the x and y coordinates of the 2012 points. As in the
previous calls specify the x coordinates with calls to rep and the y coordinates with
references to the appropriate columns of mrg.
489 CorrectAnswer: segments(rep(1, 52), mrg[, 2], rep(2, 52), mrg[, 3])
490 AnswerTests: omnitest(correctExpr='segments(rep(1, 52), mrg[, 2], rep(2, 52),
mrg[, 3])')
491 Hint: Type segments(rep(1, 52), mrg[, 2], rep(2, 52), mrg[, 3]) at the command prompt.
492
493 - Class: text
494 Output: We see from the plot that the vast majority of states have indeed improved
their particulate matter counts so the general trend is downward. There are a few
exceptions. (The topmost point in the 1999 column is actually two points that had
very close measurements.)
495
496 - Class: cmd_question
497 Output: For fun, let's see which states had higher means in 2012 than in 1999. Just
use the mrg[mrg$mean.x < mrg$mean.y, ] notation to find the rows of mrg with this
particulate property.
498 CorrectAnswer: mrg[mrg$mean.x < mrg$mean.y,]
499 AnswerTests: omnitest(correctExpr='mrg[mrg$mean.x < mrg$mean.y,]')
500 Hint: Type mrg[mrg$mean.x < mrg$mean.y,] at the command prompt.
501
502 - Class: text
503 Output: Only 4 states had worse pollution averages, and 2 of these had means that
were very close. If you want to see which states (15, 31, 35, and 40) these are, you
can check out this website https://www.epa.gov/enviro/state-fips-code-listing to
decode the state codes.
504
505 - Class: text
506 Output: This concludes the lesson, comparing air pollution data from two years in
different ways. First, we looked at measures of the entire set of monitors, then we
compared the two measures from a particular monitor, and finally, we looked at the
mean measures of the individual states.
507
508 - Class: text
509 Output: Congratulations! We hope you enjoyed this particulate lesson.
510
511 - Class: mult_question
512 Output: "Would you like to receive credit for completing this course on
Coursera.org?"
513 CorrectAnswer: NULL
514 AnswerChoices: Yes;No
515 AnswerTests: coursera_on_demand()
516 Hint: ""
517
518

```