

```

1  Course: Statistical_Inference
2  Lesson: Power
3
4  - Class: text
5  Output: "Power. (Slides for this and other Data Science courses may be found at
6  github https://github.com/DataScienceSpecialization/courses/. If you care to use
7  them, they must be downloaded as a zip file and viewed locally. This lesson
8  corresponds to 06_Statistical_Inference/11_Power.)"
9
10 - Class: text
11 Output: In this lesson, as the name suggests, we'll discuss POWER, which is the
12 probability of rejecting the null hypothesis when it is false, which is good and
13 proper.
14
15 - Class: text
16 Output: Hence you want more POWER.
17
18 - Class: text
19 Output: Power comes into play when you're designing an experiment, and in particular,
20 if you're trying to determine if a null result (failing to reject a null hypothesis)
21 is meaningful. For instance, you might have to determine if your sample size was big
22 enough to yield a meaningful, rather than random, result.
23
24 - Class: text
25 Output: Power gives you the opportunity to detect if your ALTERNATIVE hypothesis is
26 true.
27
28 - Class: mult_question
29 Output: Do you recall the definition of a Type II error? Remember, errors are bad.
30 AnswerChoices: Rejecting a true null hypothesis; Accepting a false null hypothesis;
31 Miscalculating a t score; Misspelling the word hypothesis
32 CorrectAnswer: Accepting a false null hypothesis
33 AnswerTests: omnitest(correctVal='Accepting a false null hypothesis')
34 Hint: Remember the courtroom example? Letting a guilty person walk, accepting the
35 null hypothesis of innocence, is a Type II error.
36
37 - Class: text
38 Output: Beta is the probability of a Type II error, accepting a false null
39 hypothesis; the complement of this is obviously (1 - beta) which represents the
40 probability of rejecting a false null hypothesis. This is good and this is POWER!
41
42 - Class: text
43 Output: Recall our previous example involving the Respiratory Distress Index and
44 sleep disturbances. Our null hypothesis  $H_0$  was that  $\mu = 30$  and our alternative
45 hypothesis  $H_a$  was that  $\mu > 30$ .
46
47 - Class: mult_question
48 Output: Which of the following expressions represents our test statistic under this
49 null hypothesis? Here  $\bar{X}$  represents the sample mean,  $s$  is the sample std deviation,
50 and  $n$  is the sample size. Assume  $\bar{X}$  follows a t distribution.
51 AnswerChoices:  $(\bar{X}-30)/(s/\sqrt{n})$ ;  $\bar{X}/(s^2/n)$ ;  $(\bar{X}-30)/(s^2/n)$ ;  $30/(s/\sqrt{n})$ 
52 CorrectAnswer:  $(\bar{X}-30)/(s/\sqrt{n})$ 
53 AnswerTests: omnitest(correctVal='( $\bar{X}$ -30)/(s/sqrt(n))')
54 Hint:  $(\bar{X}-30)/(s/\sqrt{n})$  measures the number of standard errors the sample mean is
55 from the mean hypothesized by  $H_0$ .
56
57 - Class: mult_question
58 Output: In the expression for the test statistic  $(\bar{X}-30)/(s/\sqrt{n})$  what does
59  $(s/\sqrt{n})$  represent?
60 AnswerChoices: a standard error; a standard bearer; a standard variance; a standard
61 sample; a standard measure
62 CorrectAnswer: a standard error
63 AnswerTests: omnitest(correctVal='a standard error')
64 Hint: Since  $(\bar{X}-30)/(s/\sqrt{n})$  measures the number of standard errors the sample
65 mean is from the mean hypothesized by  $H_0$ , the denominator is the standard error of
66 the sample mean.
67
68 - Class: text
69 Output: Suppose we're testing a null hypothesis  $H_0$  with an alpha level of .05. Since

```

H<sub>a</sub> proposes that  $\mu > 30$  (the mean hypothesized by H<sub>0</sub>), power is the probability that the true mean  $\mu$  is greater than the  $(1-\alpha)$  quantile or `qnorm(.95)`. For simplicity, assume we're working with normal distributions of which we know the variances.

```
48
49 - Class: figure
50 Output: Here's the picture we've used a lot in these lessons. As you know, the shaded
portion represents 5% of the area under the curve. If a test statistic fell in this
shaded portion we would reject H0 because the sample mean is too far from the mean
(center) of the distribution hypothesized by H0. Instead we would favor Ha, that
 $\mu > 30$ . This happens with probability .05.
51 Figure: conf_5pct.R
52 FigureType: new
53
54 - Class: text
55 Output: You might well ask, "What does this have to do with POWER?" Good question.
We'll look at some pictures to show you.
56
57 - Class: text
58 Output: First we have to emphasize a key point. The two hypotheses, H0 and Ha,
actually represent two distributions since they're talking about means or centers of
distributions. H0 says that the mean is  $\mu_0$  (30 in our example) and Ha says that
the mean is  $\mu_a$ .
59
60 - Class: text
61 Output: We're assuming normality and equal variance, say  $\sigma^2/n$ , for both
hypotheses, so under H0,  $X' \sim N(\mu_0, \sigma^2/n)$  and under Ha,  $X' \sim N(\mu_a, \sigma^2/n)$ .
62
63 - Class: figure
64 Output: Here's a picture with the two distributions. We've drawn a vertical line at
our favorite spot, at the 95th percentile of the red distribution. To the right of
the line lies 5% of the red distribution.
65 Figure: twoDistros.R
66 FigureType: new
67
68 - Class: mult_question
69 Output: Quick quiz! Which distribution represents H0?
70 AnswerChoices: the red; the blue
71 CorrectAnswer: the red
72 AnswerTests: omnitest(correctVal='the red')
73 Hint: The two distributions have the same spread (variance). They differ in their
means (centers). Which has a mean equal to that hypothesized by H0?
74
75 - Class: mult_question
76 Output: Which distribution represents Ha?
77 AnswerChoices: the red; the blue
78 CorrectAnswer: the blue
79 AnswerTests: omnitest(correctVal='the blue')
80 Hint: The two distributions have the same spread (variance). They differ in their
means (centers). Which has a mean different than that hypothesized by H0?
81
82 - Class: mult_question
83 Output: From the picture, what is the mean proposed by Ha?
84 AnswerChoices: 30; 28; 32; 36;
85 CorrectAnswer: 32
86 AnswerTests: omnitest(correctVal='32')
87 Hint: At what value of  $\mu$  does the center (highest point) of the blue
distribution fall?
88
89 - Class: text
90 Output: See how much of the blue distribution lies to the right of that big vertical
line?
91
92 - Class: text
93 Output: That, my friend, is POWER!
94
95 - Class: text
96 Output: It's the area under the blue curve (Ha) to the right of the vertical line.
97
```

```

98 - Class: figure
99 Output: Note that the placement of the vertical line depends on the null
distribution. Here's another picture with fatter distributions. The vertical line is
still at the 95th percentile of the null (red) distribution and 5% of the
distribution still lies to its right. The line is calibrated to  $\mu_0$  and the variance.
100 Figure: twoDistros_fat.R
101 FigureType: new
102
103 - Class: figure
104 Output: Back to our original picture.
105 Figure: twoDistros.R
106 FigureType: new
107
108 - Class: cmd_question
109 Output: We've shamelessly stolen plotting code from the slides so you can see  $H_a$  in
action. Let's look at pictures before we delve into numbers. We've fixed  $\mu_0$  at 30,
sigma (standard deviation) at 4 and n (sample size) at 16. The function myplot just
needs an alternative mean,  $\mu_a$ , as argument. Run myplot now with an argument of 34
to see what it does.
110 CorrectAnswer: myplot(34)
111 AnswerTests: omnitest(correctExpr='myplot(34)')
112 Hint: Type myplot(34) at the command prompt.
113
114 - Class: cmd_question
115 Output: The distribution represented by  $H_a$  moved to the right, so almost all (100%)
of the blue curve is to the right of the vertical line, indicating that with  $\mu_a=34$ ,
the test is more powerful, i.e., there's a higher probability that it's correct to
reject the null hypothesis since it appears false. Now try myplot with an argument of
33.3.
116 CorrectAnswer: myplot(33.3)
117 AnswerTests: omnitest(correctExpr='myplot(33.3)')
118 Hint: Type myplot(33.3) at the command prompt.
119
120 - Class: cmd_question
121 Output: This isn't as powerful as the test with  $\mu_a=34$  but it makes a pretty
picture. Now try myplot with an argument of 30.
122 CorrectAnswer: myplot(30)
123 AnswerTests: omnitest(correctExpr='myplot(30)')
124 Hint: Type myplot(30) at the command prompt.
125
126 - Class: text
127 Output: Uh Oh! Did the red curve disappear? No. it's just under the blue curve. The
power now, the area under the blue curve to the right of the line, is exactly 5% or
alpha!
128
129 - Class: text
130 Output: So what did we learn?
131
132 - Class: text
133 Output: First, power is a function that depends on a specific value of an alternative
mean,  $\mu_a$ , which is any value greater than  $\mu_0$ , the mean hypothesized by  $H_0$ .
(Recall that  $H_a$  specified  $\mu > 30$ .)
134
135 - Class: text
136 Output: Second, if  $\mu_a$  is much bigger than  $\mu_0=30$  then the power (probability) is
bigger than if  $\mu_a$  is close to 30. As  $\mu_a$  approaches 30, the mean under  $H_0$ , the
power approaches alpha.
137
138 - Class: cmd_question
139 Output: Just for fun try myplot with an argument of 28.
140 CorrectAnswer: myplot(28)
141 AnswerTests: omnitest(correctExpr='myplot(28)')
142 Hint: Type myplot(28) at the command prompt.
143
144 - Class: text
145 Output: We see that the blue curve has moved to the left of the red, so the area
under it, to the right of the line, is less than the 5% under the red curve. This
then is even less powerful and contradicts  $H_a$  so it's not worth looking at.

```

```

147
148 - Class: figure
149 Output: Here's a picture of the power curves for different sample sizes. Again, this
      uses code "borrowed" from the slides. The alternative means, the  $\mu_a$ 's, are plotted
      along the horizontal axis and power along the vertical.
150 Figure: powerCurve_n.R
151 FigureType: new
152
153 - Class: mult_question
154 Output: What does the graph show us about  $\mu_a$ ?
155 AnswerChoices: as it gets bigger, it gets more powerful; as it gets bigger, it gets
      less powerful; power is independent of  $\mu_a$ 
156 CorrectAnswer: as it gets bigger, it gets more powerful
157 AnswerTests: omnitest(correctVal='as it gets bigger, it gets more powerful')
158 Hint: As you move to the right along the horizontal axis, does the line go up or down?
159
160 - Class: mult_question
161 Output: What does the graph show us about sample size?
162 AnswerChoices: as it gets bigger, it gets more powerful; as it gets bigger, it gets
      less powerful; power is independent of sample size
163 CorrectAnswer: as it gets bigger, it gets more powerful
164 AnswerTests: omnitest(correctVal='as it gets bigger, it gets more powerful')
165 Hint: You have to check the color key. The purple line represents a bigger sample
      size than the red or blue. This purple line goes up faster than any of the other lines.
166
167 - Class: text
168 Output: Now back to numbers. Our test for determining rejection of  $H_0$  involved
      comparing a test statistic, namely  $Z = (X' - 30) / (\sigma / \sqrt{n})$ , against some quantile,
      say  $Z_{.95}$ , which depended on our level size alpha (.05 in this case).  $H_a$  proposed
      that  $\mu > \mu_0$ , so we tested if  $Z > Z_{.95}$ . This is equivalent to  $X' > Z_{.95} *
      (\sigma / \sqrt{n}) + 30$ , right?
169
170 - Class: text
171 Output: Recall that nifty R function pnorm, which gives us the probability that a
      value drawn from a normal distribution is greater or less than/equal to a specified
      quantile argument depending on the flag lower.tail. The function also takes a mean
      and standard deviation as arguments.
172
173 - Class: text
174 Output: Suppose we call pnorm with the quantile  $30 + Z_{.95} * (\sigma / \sqrt{n})$  and
      specify  $\mu_a$  as our mean argument. This would return a probability which we can
      interpret as POWER. Why?
175
176 - Class: figure
177 Output: Recall our picture of two distributions.  $30 + Z_{.95} * (\sigma / \sqrt{n})$ 
      represents the point at which our vertical line falls. It's the point on the null
      distribution at the  $(1-\alpha)$  level.
178 Figure: twoDistros.R
179 FigureType: new
180
181 - Class: mult_question
182 Output: Study this picture. Calling pnorm with  $30 + Z_{.95} * (\sigma / \sqrt{n})$  as the
      quantile and  $\mu_a$ , say 32, as the mean and lower.tail=FALSE does what?
183 AnswerChoices: returns the area under the blue curve to the right of the line;
      returns the area under the blue curve to the left of the line; returns the area under
      the red curve to the right of the line; returns the area under the red curve to the
      left of the line
184 CorrectAnswer: returns the area under the blue curve to the right of the line
185 AnswerTests: omnitest(correctVal='returns the area under the blue curve to the right
      of the line')
186 Hint: The argument lower.tail=FALSE is a big clue. It indicates we should look to the
      right of a line. The mean argument is the second clue. It specifies which
      distribution (red or blue) to examine.
187
188 - Class: mult_question
189 Output: Let's try some examples now. Before we do, what do we know pnorm will return
      if we specify a quantile less than the mean?
190 AnswerChoices: an answer less than .50; an answer dependent on alpha; an answer
      dependent on beta; an answer greater than 1

```

```

191 CorrectAnswer: an answer less than .50
192 AnswerTests: omnitest(correctVal='an answer less than .50')
193 Hint: There are several red herrings in the choices. First pnorm will NEVER return a
value greater than 1 because density functions by definition have areas equal to 1.
We haven't specified an alpha or beta either. The function pnorm just needs a
quantile, mean, and standard deviation. By default it looks at the lower tail of the
distribution. That leaves one choice.

194
195 - Class: cmd_question
196 Output: First, define a variable z as qnorm(.95)
197 CorrectAnswer: z <- qnorm(.95)
198 AnswerTests: expr_creates_var('z'); omnitest(correctExpr='z <- qnorm(.95)')
199 Hint: Type z <- qnorm(.95) at the command prompt.
200
201 - Class: cmd_question
202 Output: Run pnorm now with the quantile 30+z, mean=30, and lower.tail=FALSE. We've
specified sigma and n so that the standard deviation of the sample mean is 1.
203 CorrectAnswer: pnorm(30+z,mean=30,lower.tail=FALSE)
204 AnswerTests: omnitest(correctExpr='pnorm(30+z,mean=30,lower.tail=FALSE)')
205 Hint: Type pnorm(30+z,mean=30,lower.tail=FALSE) at the command prompt.
206
207 - Class: cmd_question
208 Output: That's not surprising, is it? With the mean set to mu_0 the two
distributions, null and alternative, are the same and power=alpha. Now run pnorm now
with the quantile 30+z, mean=32, and lower.tail=FALSE.
209 CorrectAnswer: pnorm(30+z,mean=32,lower.tail=FALSE)
210 AnswerTests: omnitest(correctExpr='pnorm(30+z,mean=32,lower.tail=FALSE)')
211 Hint: Type pnorm(30+z,mean=32,lower.tail=FALSE) at the command prompt.
212
213 - Class: text
214 Output: See how this is much more powerful? 64% as opposed to 5%. When the sample
mean is quite different from (many standard errors greater than) the mean
hypothesized by the null hypothesis, the probability of rejecting H_0 when it is
false is much higher. That is power!

215
216 - Class: figure
217 Output: Let's look again at the portly distributions.
218 Figure: twoDistros_fat.R
219 FigureType: new
220
221 - Class: mult_question
222 Output: With this standard deviation=2 (fatter distribution) will power be greater or
less than with the standard deviation=1?
223 AnswerChoices: greater; less than; the same
224 CorrectAnswer: less than
225 AnswerTests: omnitest(correctVal='less than')
226 Hint: A greater standard deviation means more variability in the data so the test
will be less powerful.

227
228 - Class: cmd_question
229 Output: To see this, run pnorm now with the quantile 30+z, mean=32 and sd=1. Don't
forget to set lower.tail=FALSE so you get the right tail.
230 CorrectAnswer: pnorm(30+z,mean=32,sd=1,lower.tail=FALSE)
231 AnswerTests: omnitest(correctExpr='pnorm(30+z,mean=32,sd=1,lower.tail=FALSE)')
232 Hint: Type pnorm(30+z,mean=32,sd=1,lower.tail=FALSE) at the command prompt.
233
234 - Class: cmd_question
235 Output: Now run pnorm now with the quantile 30+z*2, mean=32 and sd=2. Don't forget
to set lower.tail=FALSE so you get the right tail.
236 CorrectAnswer: pnorm(30+z*2,mean=32,sd=2,lower.tail=FALSE)
237 AnswerTests: omnitest(correctExpr='pnorm(30+z*2,mean=32,sd=2,lower.tail=FALSE)')
238 Hint: Type pnorm(30+z*2,mean=32,sd=2,lower.tail=FALSE) at the command prompt.
239
240
241 - Class: mult_question
242 Output: See the power drain from 64% to 26% ? Let's review some basic facts about
power. We saw before in our pictures that the power of the test depends on mu_a. When
H_a specifies that mu > mu_0, then as mu_a grows and exceeds mu_0 increasingly, what
happens to power?

```

```

243 AnswerChoices: it increases; it decreases; it doesn't change
244 CorrectAnswer: it increases
245 AnswerTests: omnitest(correctVal='it increases')
246 Hint: Remember the different distributions and running pnorm? As mu_a got bigger so
    did the pnorm result.
247
248 - Class: figure
249 Output: Here's another question. Recall our power curves from before.
250 Figure: powerCurve_n.R
251 FigureType: new
252
253 - Class: mult_question
254 Output: As the sample size increases, what happens to power?
255 AnswerChoices: it increases; it decreases; it doesn't change
256 CorrectAnswer: it increases
257 AnswerTests: omnitest(correctVal='it increases')
258 Hint: Look at the picture!
259
260
261 - Class: figure
262 Output: Here's another one. More power curves.
263 Figure: powerCurve_sigma.R
264 FigureType: new
265
266 - Class: mult_question
267 Output: As variance increases, what happens to power?
268 AnswerChoices: it increases; it decreases; it doesn't change
269 CorrectAnswer: it decreases
270 AnswerTests: omnitest(correctVal='it decreases')
271 Hint: Look at the picture!
272
273 - Class: figure
274 Output: Here's another one. And even more power curves.
275 Figure: powerCurve_alpha.R
276 FigureType: new
277
278 - Class: mult_question
279 Output: As alpha increases, what happens to power?
280 AnswerChoices: it increases; it decreases; it doesn't change
281 CorrectAnswer: it increases
282 AnswerTests: omnitest(correctVal='it increases')
283 Hint: Look at the picture!
284
285 - Class: text
286 Output: If  $H_a$  proposed that  $\mu \neq \mu_0$  we would calculate the one sided power using
    alpha / 2 in the direction of  $\mu_a$  (either less than or greater than  $\mu_0$ ). (This is
    only approximately right, it excludes the probability of getting a large test
    statistic in the opposite direction of the truth.
287
288 - Class: mult_question
289 Output: Since power goes up as alpha gets larger would the power of a one-sided test
    be greater or less than the power of the associated two sided test?
290 AnswerChoices: greater; less than; they're the same
291 CorrectAnswer: greater
292 AnswerTests: omnitest(correctVal='greater')
293 Hint: The quantity alpha is bigger than alpha/2 so it's got more power.
294
295 - Class: mult_question
296 Output: Finally, if  $H_a$  specified that  $\mu < \mu_0$  could we still do the same kind of
    power calculations?
297 AnswerChoices: Yes; No
298 CorrectAnswer: Yes
299 AnswerTests: omnitest(correctVal='Yes')
300 Hint: We just have to look at the right tail and flip all our reasoning.
301
302 - Class: text
303 Output: Suppose  $H_a$  says that  $\mu > \mu_0$ . Then  $\text{power} = 1 - \beta = \text{Prob} ( X' > \mu_0 +
    z_{(1-\alpha)} * \sigma/\sqrt{n})$  assuming that  $X' \sim N(\mu_a, \sigma^2/n)$ . Which quantities do
    we know in this statement, given the context of the problem? Let's work through this.

```



```

304
305 - Class: mult_question
306 Output: What does the null hypothesis  $H_0$  tell us that the population mean equals?
307 AnswerChoices: mu_0; mu_a; beta; alpha
308 CorrectAnswer: mu_0
309 AnswerTests: omnitest(correctVal='mu_0')
310 Hint:  $H_0$  simply proposes a null mean.
311
312 - Class: mult_question
313 Output: After the null mean  $\mu_0$  is proposed what does the designer of the hypothesis
test specify in order to reject or fail-to-reject  $H_0$ ? In other words, what is the
level size of the test?
314 AnswerChoices: mu_0; mu_a; beta; alpha
315 CorrectAnswer: alpha
316 AnswerTests: omnitest(correctVal='alpha')
317 Hint: The level size is indicated by a Greek letter.
318
319 - Class: text
320 Output: So we know that the quantities  $\mu_0$  and  $\alpha$  are specified by the test
designer. In the statement  $1 - \beta = \text{Prob}(X' > \mu_0 + z_{(1-\alpha)} * \sigma/\sqrt{n})$ 
given  $\mu_a > \mu_0$ ,  $\mu_0$  and  $\alpha$  are specified, and  $X'$  depends on the data. The
other four quantities, ( $\beta$ ,  $\sigma$ ,  $n$ , and  $\mu_a$ ), are all unknown.
321
322 - Class: text
323 Output: It should be obvious that specifying any three of these unknowns will allow
us to solve for the missing fourth. Usually, you only try to solve for power ( $1-\beta$ )
or the sample size  $n$ .
324
325 - Class: text
326 Output: An interesting point is that power doesn't need  $\mu_a$ ,  $\sigma$  and  $n$ 
individually. Instead only  $\sqrt{n} * (\mu_a - \mu_0) / \sigma$  is needed. The quantity
 $(\mu_a - \mu_0) / \sigma$  is called the EFFECT SIZE. This is the difference in the means
in standard deviation units. It is unit free so it can be interpreted in different
settings.
327
328 - Class: text
329 Output: We'll work through some examples of this now. However, instead of assuming
that we're working with normal distributions let's work with  $t$  distributions.
Remember, they're pretty close to normal with large enough sample sizes.
330
331 - Class: text
332 Output: Power is still a probability, namely  $P((X' - \mu_0)/(S/\sqrt{n}) >
t_{(1-\alpha, n-1)} \text{ given } H_a \text{ that } \mu > \mu_a)$ . Notice we use the  $t$  quantile instead of
the  $z$ . Also, since the proposed distribution is not centered at  $\mu_0$ , we have to use
the non-central  $t$  distribution.
333
334 - Class: text
335 Output: R comes to the rescue again with the function power.t.test. We can omit one
of the arguments and the function solves for it. Let's first use it to solve for power.
336
337 - Class: text
338 Output: We'll run it three times with the same values for  $n$  (16) and  $\alpha$  (.05) but
different delta and standard deviation values. We'll show that if delta (difference
in means) divided by the standard deviation is the same, the power returned will also
be the same. In other words, the effect size is constant for all three of our tests.
339
340 - Class: cmd_question
341 Output: We'll specify a positive delta; this tells power.t.test that  $H_a$  proposes
that  $\mu > \mu_0$  and so we'll need a one-sided test. First run power.t.test(n = 16,
delta = 2 / 4, sd=1, type = "one.sample", alt = "one.sided")$power .
342 CorrectAnswer: power.t.test(n = 16, delta = 2 / 4, sd=1, type = "one.sample", alt =
"one.sided")$power
343 AnswerTests: omnitest(correctExpr='power.t.test(n = 16, delta = 2 / 4, sd=1, type =
"one.sample", alt = "one.sided")$power')
344 Hint: Type power.t.test(n = 16, delta = 2 / 4, sd=1, type = "one.sample", alt =
"one.sided")$power at the prompt.
345
346 - Class: cmd_question
347 Output: Now change delta to 2 and sd to 4. Keep everything else the same.

```

```

348 CorrectAnswer: power.t.test(n = 16, delta = 2 , sd=4, type = "one.sample", alt =
"one.sided")$power
349 AnswerTests: omnitest(correctExpr='power.t.test(n = 16, delta = 2 , sd=4, type =
"one.sample", alt = "one.sided")$power')
350 Hint: Type power.t.test(n = 16, delta = 2 , sd=4, type = "one.sample", alt =
"one.sided")$power at the prompt.
351
352 - Class: cmd_question
353 Output: Same answer, right? Now change delta to 100 and sd to 200. Keep everything
else the same.
354 CorrectAnswer: power.t.test(n = 16, delta = 100 , sd=200, type = "one.sample", alt =
"one.sided")$power
355 AnswerTests: omnitest(correctExpr='power.t.test(n = 16, delta = 100 , sd=200, type =
"one.sample", alt = "one.sided")$power')
356 Hint: Type power.t.test(n = 16, delta = 100 , sd=200, type = "one.sample", alt =
"one.sided")$power at the prompt.
357
358 - Class: text
359 Output: So keeping the effect size (the ratio delta/sd) constant preserved the power.
Let's try a similar experiment except now we'll specify a power we want and solve for
the sample size n.
360
361 - Class: cmd_question
362 Output: First run power.t.test(power = .8, delta = 2 / 4, sd=1, type =
"one.sample", alt = "one.sided")$n .
363 CorrectAnswer: power.t.test(power = .8, delta = 2 / 4, sd=1, type = "one.sample",
alt = "one.sided")$n
364 AnswerTests: omnitest(correctExpr='power.t.test(power = .8, delta = 2 / 4, sd=1, type
= "one.sample", alt = "one.sided")$n')
365 Hint: Type power.t.test(power = .8, delta = 2 / 4, sd=1, type = "one.sample", alt =
"one.sided")$n.
366
367 - Class: cmd_question
368 Output: Now change delta to 2 and sd to 4. Keep everything else the same.
369 CorrectAnswer: power.t.test(power = .8, delta = 2, sd=4, type = "one.sample", alt =
"one.sided")$n
370 AnswerTests: omnitest(correctExpr='power.t.test(power = .8, delta = 2, sd=4, type =
"one.sample", alt = "one.sided")$n')
371 Hint: Type power.t.test(power = .8, delta = 2, sd=4, type = "one.sample", alt =
"one.sided")$n.
372
373 - Class: cmd_question
374 Output: Same answer, right? Now change delta to 100 and sd to 200. Keep everything
else the same.
375 CorrectAnswer: power.t.test(power = .8, delta = 100 , sd=200, type = "one.sample",
alt = "one.sided")$n
376 AnswerTests: omnitest(correctExpr='power.t.test(power = .8, delta = 100 , sd=200,
type = "one.sample", alt = "one.sided")$n')
377 Hint: Type power.t.test(power = .8, delta = 100 , sd=200, type = "one.sample", alt =
"one.sided")$n.
378
379 - Class: cmd_question
380 Output: Now use power.t.test to find delta for a power=.8 and n=26 and sd=1
381 CorrectAnswer: power.t.test(power = .8, n=26, sd=1, type = "one.sample", alt =
"one.sided")$delta
382 AnswerTests: omnitest(correctExpr='power.t.test(power = .8, n=26, sd=1, type =
"one.sample", alt = "one.sided")$delta')
383 Hint: Type power.t.test(power = .8, n=26, sd=1, type = "one.sample", alt =
"one.sided")$delta.
384
385 - Class: cmd_question
386 Output: Not a surprising result, is it? It told you before that with an effect size
of .5 and power .8, you need a sample size a little more than 26. Now run it with n=27.
387 CorrectAnswer: power.t.test(power = .8, n=27, sd=1, type = "one.sample", alt =
"one.sided")$delta
388 AnswerTests: omnitest(correctExpr='power.t.test(power = .8, n=27, sd=1, type =
"one.sample", alt = "one.sided")$delta')
389 Hint: Type power.t.test(power = .8, n=27, sd=1, type = "one.sample", alt =
"one.sided")$delta.

```



```

390
391 - Class: mult_question
392 Output: What do you think will happen if you doubled sd to 2 and ran the same test?
393 AnswerChoices: delta will double; delta will halve; delta won't change
394 CorrectAnswer: delta will double
395 AnswerTests: omnitest(correctVal='delta will double')
396 Hint: Since you're doubling the denominator (sd) you have to double the numerator
    (delta) in order to keep the effect size constant.
397
398 - Class: text
399 Output: Now for a quick review. We call this the power.u.test since it comes after
    the power.t.test. LOL.
400
401 - Class: mult_question
402 Output: 1. The level of a test is specified by what?
403 AnswerChoices: alpha; beta; gamma; delta; None of the others
404 CorrectAnswer: alpha
405 AnswerTests: omnitest(correctVal='alpha')
406 Hint: The level refers to the probability of rejecting the null hypothesis when it
    is true. This is the first Greek letter.
407
408 - Class: mult_question
409 Output: 2. What is a Type II error?
410 AnswerChoices: rejecting a true hypothesis; rejecting a false hypothesis;
    accepting a true hypothesis; accepting a false hypothesis
411 CorrectAnswer: accepting a false hypothesis
412 AnswerTests: omnitest(correctVal='accepting a false hypothesis')
413 Hint: Only two of the choices are errors. Eliminate those and then eliminate the Type
    I error.
414
415 - Class: mult_question
416 Output: 3. What is power?
417 AnswerChoices: thrilling; alpha; beta; gamma; delta; None of the others
418 CorrectAnswer: None of the others
419 AnswerTests: omnitest(correctVal='None of the others')
420 Hint: Power is 1-beta. Is that in the list?
421
422 - Class: mult_question
423 Output: 4. You're a perfectionist designing an experiment and you want both alpha and
    beta to be small. Can they both be 0 for this single test?
424 AnswerChoices: Yes; No
425 CorrectAnswer: No
426 AnswerTests: omnitest(correctVal='No')
427 Hint: If alpha=0, then you ALWAYS fail to reject (that is, you accept)  $H_0$ ; if  $H_0$  is
    false, beta=1. Similarly if beta=0, you ALWAYS reject  $H_0$ ; if  $H_0$  is true then
    alpha=1.
428
429 - Class: mult_question
430 Output: 5. Suppose  $H_0$  proposes  $\mu = \mu_0$  and  $H_a$  proposes that  $\mu < \mu_0$ . You'll
    test a series of  $\mu_a$  with power  $\neq \alpha$ . Which of the following is NOT true?
431 AnswerChoices:  $\mu_0 - \mu_a > 0$ ;  $\mu_a - \mu_0 < 0$ ;  $\mu_a - \mu_0 = 0$ ; huh?
432 CorrectAnswer:  $\mu_a - \mu_0 = 0$ 
433 AnswerTests: omnitest(correctVal='mu_a-mu_0=0')
434 Hint: As in the case in which  $H_a$  proposes that  $\mu > \mu_0$ , when  $\mu_a = \mu_0$ ,
    alpha=power. This was excluded in the problem statement so  $\mu_a = \mu_0$  cannot be true
    in this case.
435
436 - Class: mult_question
437 Output: 6. Suppose  $H_0$  proposes  $\mu = \mu_0$  and  $H_a$  proposes that  $\mu < \mu_0$ . Which of
    the following is true?
438 AnswerChoices: the smaller  $\mu_0 - \mu_a$  the more powerful the test; the smaller
     $\mu_a - \mu_0$  the more powerful the test;  $\mu_0 = \mu_a$  maximizes the power
439 CorrectAnswer: the smaller  $\mu_a - \mu_0$  the more powerful the test
440 AnswerTests: omnitest(correctVal='the smaller mu_a-mu_0 the more powerful the test')
441 Hint: Here  $\mu_a < \mu_0$  and the smaller  $\mu_a - \mu_0$  is, the easier it is to discriminate
    between  $\mu_a$  and  $\mu_0$ .
442
443 - Class: mult_question
444 Output: 7. Which expression represents the size effect?

```

```

445 AnswerChoices: (mu_a - mu_0) / sigma; (mu_a - mu_0) / n; (mu_a - mu_0) / sqrt(sigma);
    (mu_a - mu_0) / sqrt(n)
446 CorrectAnswer: (mu_a - mu_0) / sigma
447 AnswerTests: omnitest(correctVal='(mu_a - mu_0) / sigma')
448 Hint: The size effect is the distance between the two proposed means in standard
    deviation units.
449
450 - Class: mult_question
451 Output: 8. True or False? More power is better than less power.
452 AnswerChoices: True; False
453 CorrectAnswer: True
454 AnswerTests: omnitest(correctVal='True')
455 Hint: Power is good.
456
457 - Class: mult_question
458 Output: 9. True or False? A larger beta (call it beta_max) is more powerful than a
    smaller beta.
459 AnswerChoices: True; False
460 CorrectAnswer: False
461 AnswerTests: omnitest(correctVal='False')
462 Hint: Recall that beta is the probability of an error (Type II), so you'd like a
    smaller beta.
463
464 - Class: mult_question
465 Output: 10. True or False? The larger the sample size the less powerful the test.
466 AnswerChoices: True; False
467 CorrectAnswer: False
468 AnswerTests: omnitest(correctVal='False')
469 Hint: It's usually better to have more data to work with.
470
471 - Class: text
472 Output: Congrats! You finished this powerful lesson. We hope you feel emPOWERED.
473
474 - Class: mult_question
475 Output: "Would you like to receive credit for completing this course on
476 Coursera.org?"
477 CorrectAnswer: NULL
478 AnswerChoices: Yes;No
479 AnswerTests: coursera_on_demand()
480 Hint: ""
481

```